University of BRISTOL

Machine Learning

Linear Regression

Carl Henrik Ek - carlhenrik.ek@bristol.ac.uk

October 9, 2018

http://www.carlhenrik.com

Introduction

So Far

- Lecture 1 What is machine Learning
  - assumptions are the fundation of learning
  - probabilities are the language of assumptions

So Far

- Lecture 1 What is machine Learning
  - assumptions are the fundation of learning
  - probabilities are the language of assumptions
- Lecture 2 Probabilities
  - what are the rules of probability
  - distributions are the parametrised form of a probability

So Far

- Lecture 1 What is machine Learning
  - assumptions are the fundation of learning
  - probabilities are the language of assumptions
- Lecture 2 Probabilities
  - what are the rules of probability
  - distributions are the parametrised form of a probability
- Lecture 3 Distributions
  - discrete and continous distributions
  - conjugate distributions

Conjugacy

$$= P(Y|\theta) \cdot p(\theta) \frac{1}{\int p(\mathbf{Y}|\theta) p(\theta) \mathrm{d}\theta} \qquad \propto P(Y|\theta)) \cdot p(\theta)$$

- A conjugate prior to a likelihood is such that the prior and the posterior is in the same functional family
- Knowing the form of the posterior allows us to avoid computing the evidence and just identify parameters

Linear Regression

Linear Regression [1] Ch 3.1

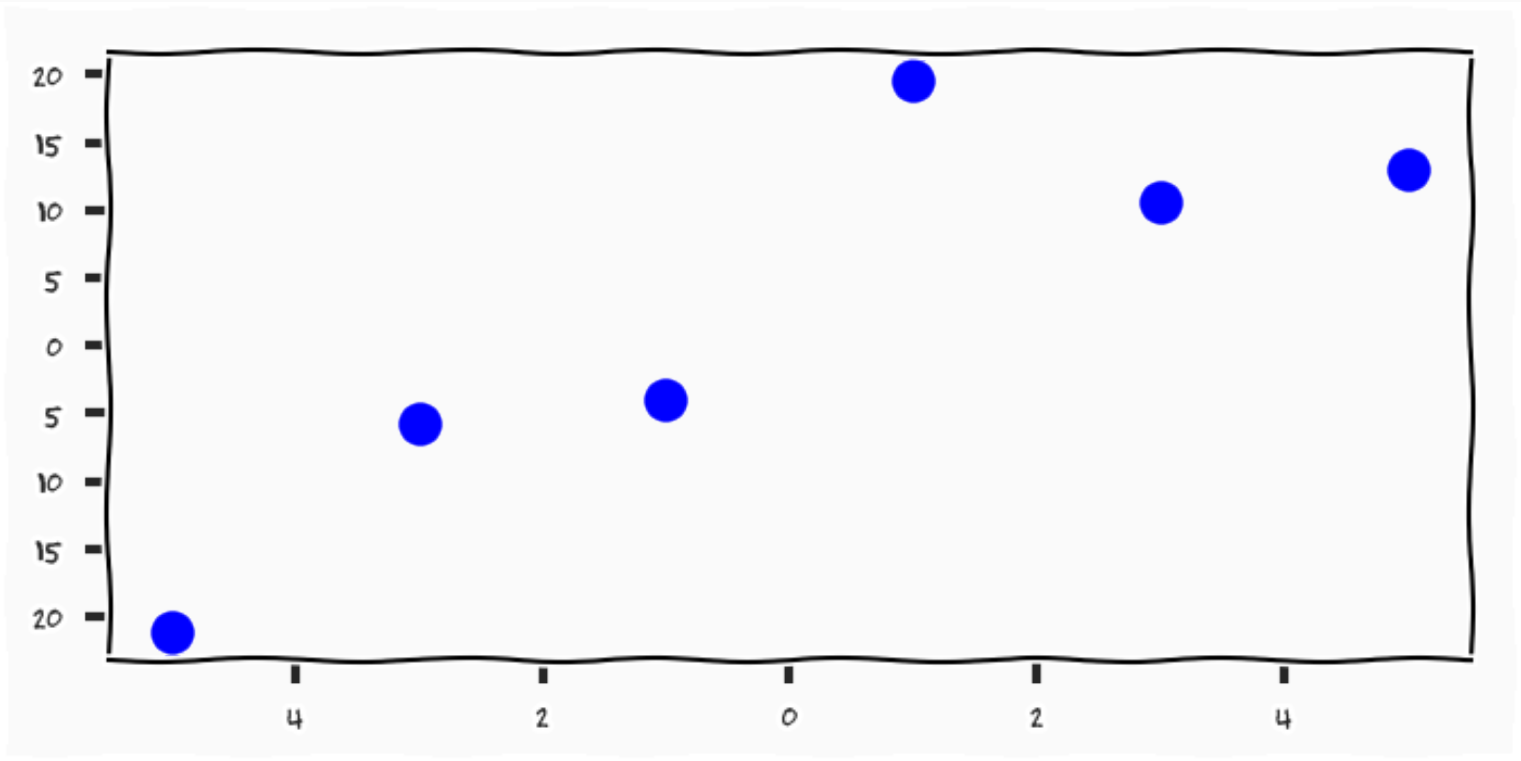

- Linear function in both parameters and data

$$y(\mathbf{x}, \mathbf{w}) = w_0 + w_1 x_1 + \ldots w_D x_D = \mathbf{w}^{\mathrm{T}}\mathbf{x} + w_0 = \{D = 1\} w_0 + w_1 * x$$

Linear Regression

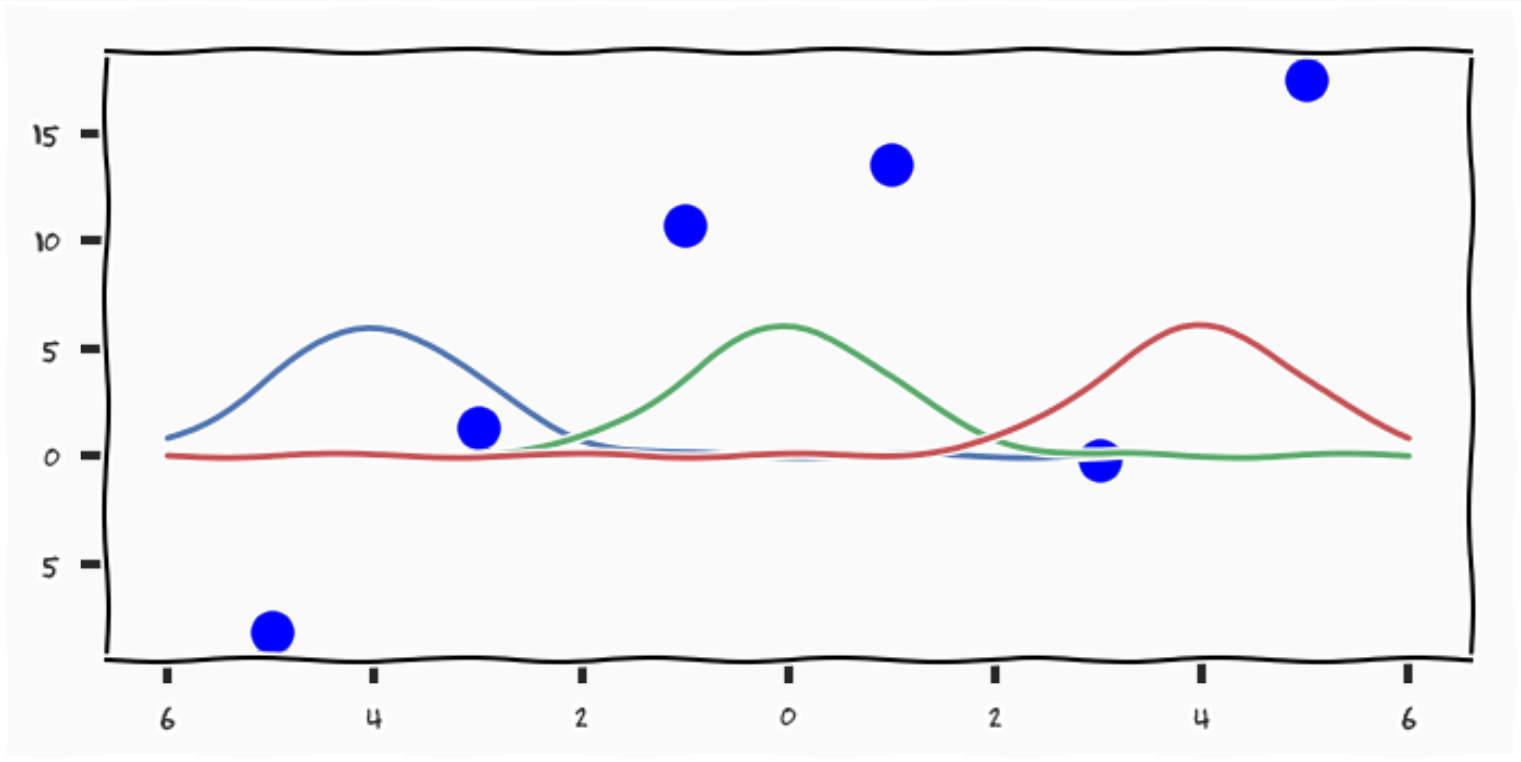

- Linear function only in parameters

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x}) = \{\phi_0(\mathbf{x}) = 1\} = \mathbf{w}^{\mathrm{T}} \boldsymbol{\phi}(\mathbf{x})$$

- We can choose many types of basis functions $\phi(\mathbf{x})$

Linear Regression

Model

$$t = y(\mathbf{x}, \mathbf{w}) + \epsilon = \mathbf{w}^{\mathrm{T}} \boldsymbol{\phi}(\mathbf{x}) + \epsilon$$
$$\epsilon \sim \mathcal{N}(0, \beta^{-1} I)$$

- We assume that we have been given data pairs $\{t_i, \mathbf{x}\}_{i=1}^{N}$ corrupted by addative noise
- We assume that the distribution of the noise follows a Gaussian

Linear Regression

Task 1 define a likelihood

- what output do I consider likely under a given model?

Task 2 define an assumption of the model

- what types of models do I think are more probable than others
- $\Rightarrow$ what are my beliefs, i.e formulate prior

Task 3 update my belief with new observations

- formulate posterior

Task 4 predict using my new belief

- formulate predictive distribution

Likelihood

$$t = \mathbf{w}^{\mathrm{T}} \boldsymbol{\phi}(\mathbf{x}) + \epsilon$$

Likelihood

$$t = \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) + \epsilon$$

$$t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) = \epsilon$$

Likelihood

$$t = \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) + \epsilon$$

$$t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) = \epsilon$$

$$t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) \sim \mathcal{N}(\epsilon|0, \beta^{-1}I) = \left(\frac{\beta}{2\pi}\right)^{\frac{1}{2}} e^{-\frac{1}{2}(\epsilon-0)\beta(\epsilon-0)}$$

Likelihood

$$t = \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) + \epsilon$$

$$t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) = \epsilon$$

$$t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) \sim \mathcal{N}(\epsilon|0, \beta^{-1}I) = \left(\frac{\beta}{2\pi}\right)^{\frac{1}{2}} e^{-\frac{1}{2}(\epsilon-0)\beta(\epsilon-0)}$$

$$\Rightarrow \mathcal{N}(t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}))|0, \beta^{-1}I) = \left(\frac{\beta}{2\pi}\right)^{\frac{1}{2}} e^{-\frac{1}{2}(t-\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}))\beta(t-\mathbf{w}^{\mathrm{T}}}$$

Likelihood

$$t = \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) + \epsilon$$

$$t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) = \epsilon$$

$$t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) \sim \mathcal{N}(\epsilon|0, \beta^{-1}I) = \left(\frac{\beta}{2\pi}\right)^{\frac{1}{2}} e^{-\frac{1}{2}(\epsilon-0)\beta(\epsilon-0)}$$

$$\Rightarrow \mathcal{N}(t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}))|0, \beta^{-1}I) = \left(\frac{\beta}{2\pi}\right)^{\frac{1}{2}} e^{-\frac{1}{2}(t-\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}))\beta(t-\mathbf{w}^{\mathrm{T}}$$

$$\Rightarrow \mathcal{N}(t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}))|0, \beta^{-1}I) = \mathcal{N}(t|\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x})), \beta^{-1}I)$$

Likelihood

$$t = \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) + \epsilon$$

$$t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) = \epsilon$$

$$t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) \sim \mathcal{N}(\epsilon|0, \beta^{-1}I) = \left(\frac{\beta}{2\pi}\right)^{\frac{1}{2}} e^{-\frac{1}{2}(\epsilon-0)\beta(\epsilon-0)}$$

$$\Rightarrow \mathcal{N}(t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}))|0, \beta^{-1}I) = \left(\frac{\beta}{2\pi}\right)^{\frac{1}{2}} e^{-\frac{1}{2}(t-\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}))\beta(t-\mathbf{w}^{\mathrm{T}}$$

$$\Rightarrow \mathcal{N}(t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}))|0, \beta^{-1}I) = \mathcal{N}(t|\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x})), \beta^{-1}I)$$

$$\Rightarrow p(t|\mathbf{w}, \mathbf{x}) = \mathcal{N}(t|\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x})), \beta^{-1}I)$$

Likelihood

$$t = \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) + \epsilon$$

$$t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) = \epsilon$$

$$t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}) \sim \mathcal{N}(\epsilon|0, \beta^{-1}I) = \left(\frac{\beta}{2\pi}\right)^{\frac{1}{2}} e^{-\frac{1}{2}(\epsilon-0)\beta(\epsilon-0)}$$

$$\Rightarrow \mathcal{N}(t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}))|0, \beta^{-1}I) = \left(\frac{\beta}{2\pi}\right)^{\frac{1}{2}} e^{-\frac{1}{2}(t-\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}))\beta(t-\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}))}$$

$$\Rightarrow \mathcal{N}(t - \mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}))|0, \beta^{-1}I) = \mathcal{N}(t|\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x})), \beta^{-1}I)$$

$$\Rightarrow p(t|\mathbf{w}, \mathbf{x}) = \mathcal{N}(t|\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x})), \beta^{-1}I)$$

- By making an assumption of the noise we have reached a conditional distribution over the output given the model, i.e.

Likelihood

- Likelihood

$$p(t|\mathbf{x}, \mathbf{w}, \beta) = \mathcal{N}\left(t|\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}), \beta^{-1}\right)$$

- Independence

$$p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \prod_{n=1}^{N} \mathcal{N}\left(t|\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}), \beta^{-1}\right)$$

Assume each output to be independent given the input and the parameters

Maximum Likelihood

- If we want we can avoid using our belief and simply pick the model that maximises our likelihood
- In this setting you can think of the likelihood as a quantification of an error
- Find the parameters that minimises the error

Maximum Likelihood

- If we want we can avoid using our belief and simply pick the model that maximises our likelihood
- In this setting you can think of the likelihood as a quantification of an error
- Find the parameters that minimises the error
- *Why is this a scary thing to do?*

Maximum Likelihood [1] Ch 3.1.1

$$p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \prod_{n=1}^{N} \mathcal{N}\left(t|\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}_n), \beta^{-1}\right)$$

Maximum Likelihood [1] Ch 3.1.1

$$p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \prod_{n=1}^{N} \mathcal{N}\left(t|\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}_n), \beta^{-1}\right)$$

$$= \prod_{n=1}^{N} \frac{1}{(2\pi\beta^{-1})^{\frac{1}{2}}} e^{-\frac{1}{2}\beta(t_n - \mathbf{w}^{\mathrm{T}}\phi(\mathbf{x}_n))^2}$$

Maximum Likelihood [1] Ch 3.1.1

$$
\begin{aligned}
p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) &= \prod_{n=1}^{N} \mathcal{N}\left(t|\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}_n), \beta^{-1}\right) \\
&= \prod_{n=1}^{N} \frac{1}{(2\pi\beta^{-1})^{\frac{1}{2}}} e^{-\frac{1}{2}\beta(t_n - \mathbf{w}^{\mathrm{T}}\phi(\mathbf{x}_n))^2} \\
&= \left(\frac{\beta}{2\pi}\right)^{\frac{N}{2}} e^{-\frac{\beta}{2}\sum_{n=1}^{N}(t_n - \mathbf{w}^{\mathrm{T}}\phi(\mathbf{x}_n))^2}
\end{aligned}
$$

Maximum Likelihood [1] Ch 3.1.1

$$
\begin{aligned}
p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) &= \prod_{n=1}^{N} \mathcal{N}\left(t|\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}(\mathbf{x}_n), \beta^{-1}\right) \\
&= \prod_{n=1}^{N} \frac{1}{(2\pi\beta^{-1})^{\frac{1}{2}}} e^{-\frac{1}{2}\beta(t_n - \mathbf{w}^{\mathrm{T}}\phi(\mathbf{x}_n))^2} \\
&= (\frac{\beta}{2\pi})^{\frac{N}{2}} e^{-\frac{\beta}{2}\sum_{n=1}^{N}(t_n - \mathbf{w}^{\mathrm{T}}\phi(\mathbf{x}_n))^2} \\
\log p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) &= \frac{N}{2}(\log(\beta) - \log(2\pi)) - \beta\frac{1}{2}\sum_{n=1}^{N}(t_n - \mathbf{w}^{\mathrm{T}}\phi(\mathbf{x}_n))^2
\end{aligned}
$$

Maximum Likelihood

$$\log p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \frac{N}{2}(\underbrace{\log(\beta)}_{\text{A}} - \underbrace{\log(2\pi)}_{\text{B}}) - \underbrace{\beta\frac{1}{2}\sum_{n=1}^{N}(t_n - \mathbf{w}^{\mathrm{T}}\phi(\mathbf{x}_n))^2}_{\text{C}}$$

A noise precision

B constant

C error

Maximum Likelihood

- Take derivative

$$\nabla \log\, p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \beta \sum_{n=1}^{N} (t_n - \mathbf{w}^{\mathrm{T}} \phi(\mathbf{x}_n)) \phi(\mathbf{x}_n)^{\mathrm{T}}$$

Maximum Likelihood

- Take derivative

$$\nabla \log\, p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \beta \sum_{n=1}^{N} (t_n - \mathbf{w}^{\mathrm{T}} \phi(\mathbf{x}_n)) \phi(\mathbf{x}_n)^{\mathrm{T}}$$

- Stationary point

$$0 = \sum_{n=1}^{N} t_n \phi(\mathbf{x}_n)^{\mathrm{T}} - \mathbf{w}^{\mathrm{T}} \left(\sum_{n=1}^{N} \phi(\mathbf{x}_n) \phi(\mathbf{x}_n)^{\mathrm{T}} \right)$$

Maximum Likelihood

- Take derivative

$$\nabla \log\, p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \beta \sum_{n=1}^{N} (t_n - \mathbf{w}^{\mathrm{T}} \phi(\mathbf{x}_n)) \phi(\mathbf{x}_n)^{\mathrm{T}}$$

- Stationary point

$$0 = \sum_{n=1}^{N} t_n \phi(\mathbf{x}_n)^{\mathrm{T}} - \mathbf{w}^{\mathrm{T}} \left(\sum_{n=1}^{N} \phi(\mathbf{x}_n) \phi(\mathbf{x}_n)^{\mathrm{T}} \right)$$

- Solve for parameters $\mathbf{w}$

$$\mathbf{w}_{\mathrm{ML}} = (\phi(\mathbf{X})^{\mathrm{T}} \phi(\mathbf{X}))^{-1} \phi(\mathbf{X})^{\mathrm{T}} \mathbf{t}$$

Maximum Likelihood

- Take derivative

$$\nabla \log\, p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \beta \sum_{n=1}^{N} (t_n - \mathbf{w}^{\mathrm{T}} \phi(\mathbf{x}_n)) \phi(\mathbf{x}_n)^{\mathrm{T}}$$

- Stationary point

$$0 = \sum_{n=1}^{N} t_n \phi(\mathbf{x}_n)^{\mathrm{T}} - \mathbf{w}^{\mathrm{T}} \left(\sum_{n=1}^{N} \phi(\mathbf{x}_n) \phi(\mathbf{x}_n)^{\mathrm{T}} \right)$$

- Solve for parameters $\mathbf{w}$

$$\mathbf{w}_{\mathrm{ML}} = (\phi(\mathbf{X})^{\mathrm{T}} \phi(\mathbf{X}))^{-1} \phi(\mathbf{X})^{\mathrm{T}} \mathbf{t}$$

- and precision

$$\frac{1}{\beta_{\mathrm{ML}}} = \frac{1}{N} \sum_{n=1}^{N} (t_n - \mathbf{w}_{\mathrm{ML}}^{\mathrm{T}} \phi(\mathbf{x}_n))^2$$

Maximum Likelihood

$$\mathbf{w}_{\mathsf{ML}} = \underbrace{(\phi(\mathbf{X})^{\mathrm{T}}\phi(\mathbf{X}))^{-1}\phi(\mathbf{X})^{\mathrm{T}}}_{\phi(\mathbf{X})^{+}}\mathbf{t}$$

- Moore-Penrose inverse (`np.linalg.pinv` in numpy)

Bayesian Linear Regression [1] Ch 3.3

- Likelihood is Gaussian in $\mathbf{w}$

Bayesian Linear Regression [1] Ch 3.3

- Likelihood is Gaussian in $\mathbf{w}$
- Conjugate Prior

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_0, \mathbf{S}_0)$$

Bayesian Linear Regression [1] Ch 3.3

- Likelihood is Gaussian in $\mathbf{w}$
- Conjugate Prior

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_0, \mathbf{S}_0)$$

- Posterior

$$p(w|\mathbf{t}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_N, \mathbf{S}_N)$$

Bayesian Linear Regression [1] Ch 3.3

- Likelihood is Gaussian in $\mathbf{w}$
- Conjugate Prior

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_0, \mathbf{S}_0)$$

- Posterior

$$p(w|\mathbf{t}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_N, \mathbf{S}_N)$$

- $\mathbf{m}_N, \mathbf{S}_N$ is the mean and the co-variance of the posterior after having seen N data-points

Bayesian Linear Regression [1] Ch 3.3

- Likelihood is Gaussian in $\mathbf{w}$
- Conjugate Prior

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_0, \mathbf{S}_0)$$

- Posterior

$$p(w|\mathbf{t}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_N, \mathbf{S}_N)$$

- $\mathbf{m}_N, \mathbf{S}_N$ is the mean and the co-variance of the posterior after having seen N data-points
- Gaussian identities

Posterior

- Posterior is Gaussian

$$p(\mathbf{w}|\mathbf{t}, \mathbf{X}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_N, \mathbf{S}_N)$$

Posterior

- Posterior is Gaussian

$$p(\mathbf{w}|\mathbf{t}, \mathbf{X}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_N, \mathbf{S}_N)$$

- Identification

$$p(\mathbf{w}|\mathbf{t}, \mathbf{X}) \propto p(\mathbf{t}|\mathbf{X}, \mathbf{w})p(\mathbf{w})$$

Posterior

- Posterior is Gaussian

$$p(\mathbf{w}|\mathbf{t}, \mathbf{X}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_N, \mathbf{S}_N)$$

- Identification

$$p(\mathbf{w}|\mathbf{t}, \mathbf{X}) \propto p(\mathbf{t}|\mathbf{X}, \mathbf{w})p(\mathbf{w})$$

- Posterior

$$\mathbf{m}_N = \left(\mathbf{S}_0^{-1} + \beta\phi(\mathbf{X})^{\mathrm{T}}\phi(\mathbf{X})\right)^{-1}\left(S_0^{-1}\mathbf{m}_0 + \beta\phi(\mathbf{X})^{\mathrm{T}}\mathbf{t}\right)$$
$$\mathbf{S}_N = \left(\mathbf{S}_0^{-1} + \beta\phi(\mathbf{X})^{\mathrm{T}}\phi(\mathbf{X})\right)^{-1}$$

Posterior

- Assumption Zero mean isotropic Gaussian

$$p(\mathbf{w}|\alpha) = \mathcal{N}(\mathbf{w}|0, \alpha^{-1}\mathbf{I})$$

Posterior

- Assumption Zero mean isotropic Gaussian

$$p(\mathbf{w}|\alpha) = \mathcal{N}(\mathbf{w}|0, \alpha^{-1}\mathbf{I})$$

- Posterior

$$p(\mathbf{w}|\mathbf{t}, \mathbf{X}) = \mathcal{N}(\mathbf{w}|\beta \left(\alpha\mathbf{I} + \beta\phi(\mathbf{X})^{\mathrm{T}}\phi(\mathbf{X})\right)^{-1} \phi(\mathbf{X})^{\mathrm{T}}\mathbf{t},$$
$$\left(\alpha\mathbf{I} + \beta\phi(\mathbf{X})^{\mathrm{T}}\phi(\mathbf{X})\right)^{-1})$$

Posterior

- Assumption Zero mean isotropic Gaussian

$$p(\mathbf{w}|\alpha) = \mathcal{N}(\mathbf{w}|0, \alpha^{-1}\mathbf{I})$$

- Posterior

$$p(\mathbf{w}|\mathbf{t}, \mathbf{X}) = \mathcal{N}(\mathbf{w}|\beta \left(\alpha\mathbf{I} + \beta\phi(\mathbf{X})^{\mathrm{T}}\phi(\mathbf{X})\right)^{-1} \phi(\mathbf{X})^{\mathrm{T}}\mathbf{t},$$
$$\left(\alpha\mathbf{I} + \beta\phi(\mathbf{X})^{\mathrm{T}}\phi(\mathbf{X})\right)^{-1})$$

- ML

$$\mathbf{w}_{\mathrm{ML}} = (\phi(\mathbf{X})^{\mathrm{T}}\phi(\mathbf{X}))^{-1}\phi(\mathbf{X})^{\mathrm{T}}\mathbf{t}$$

Linear Regression Example [1] Figure 3.7

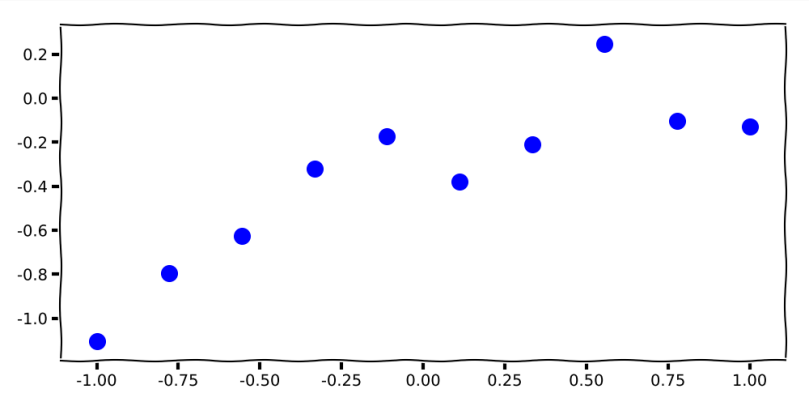

- Model

$$y(x, \mathbf{w}) = w_0 + w_1 x$$

- Data

$$f(x, \mathbf{a}) = a_0 + a_1 x, \ \{a_0, a_1\} = \{-0.3, 0.5\}$$

$$t = f(x, \mathbf{a}) + \epsilon, \ \epsilon \sim \mathcal{N}(0, 0.2^2)$$

- Prior

$$p(\mathbf{w}) = \mathcal{N}(\boldsymbol{w}|\mathbf{0}, 2.0 \cdot \mathbf{I})$$

Linear Regression Example

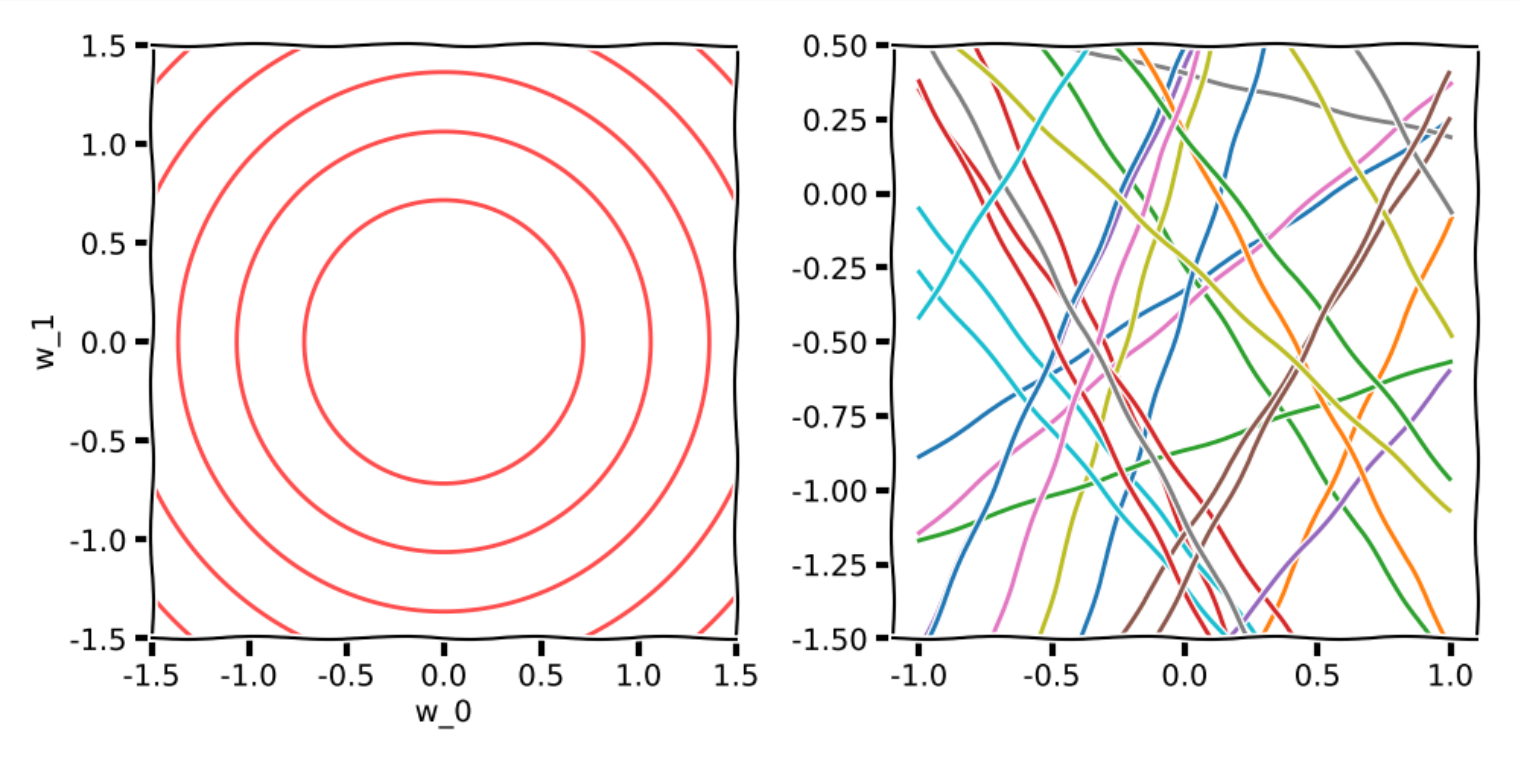

Linear Regression Example

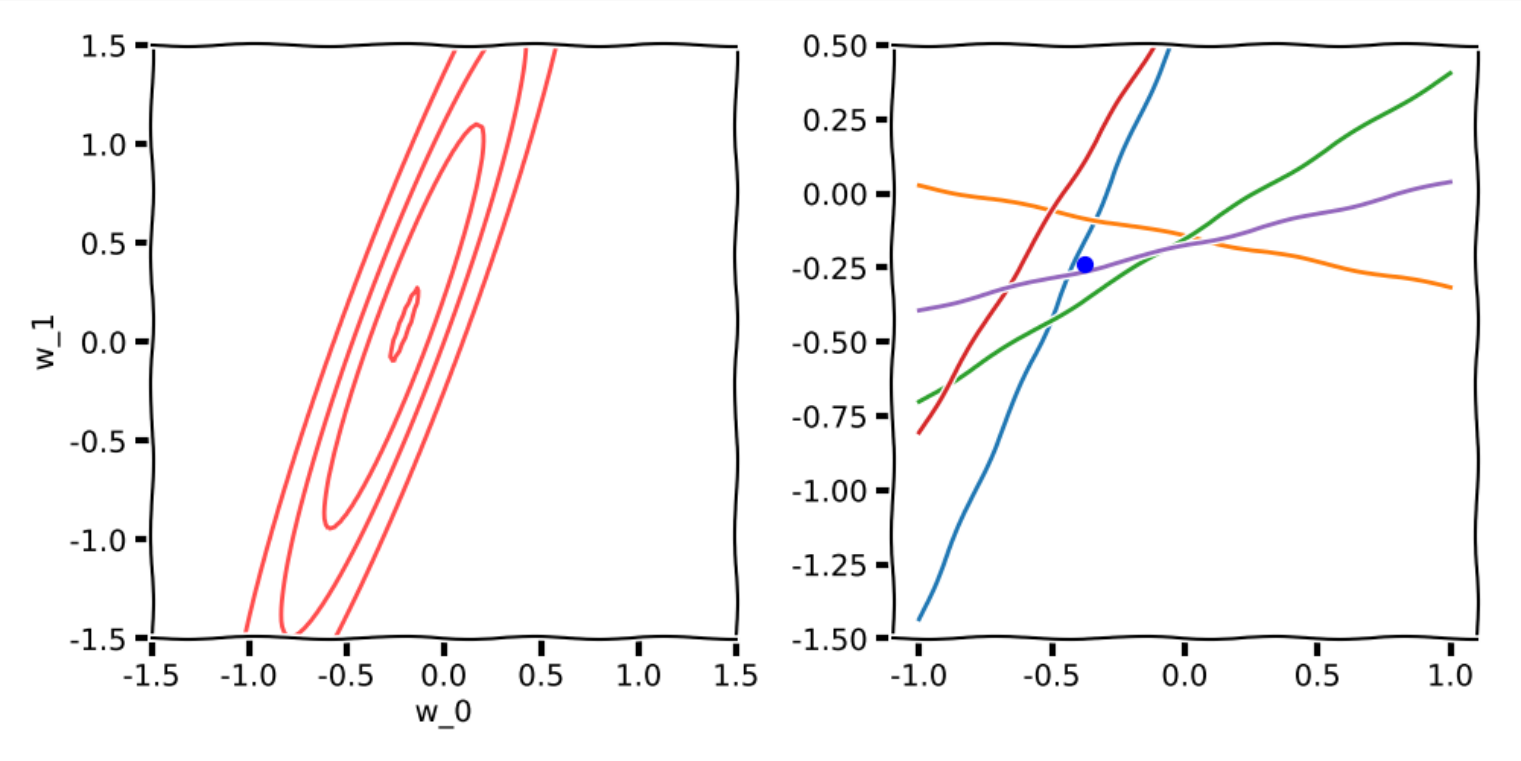

Linear Regression Example

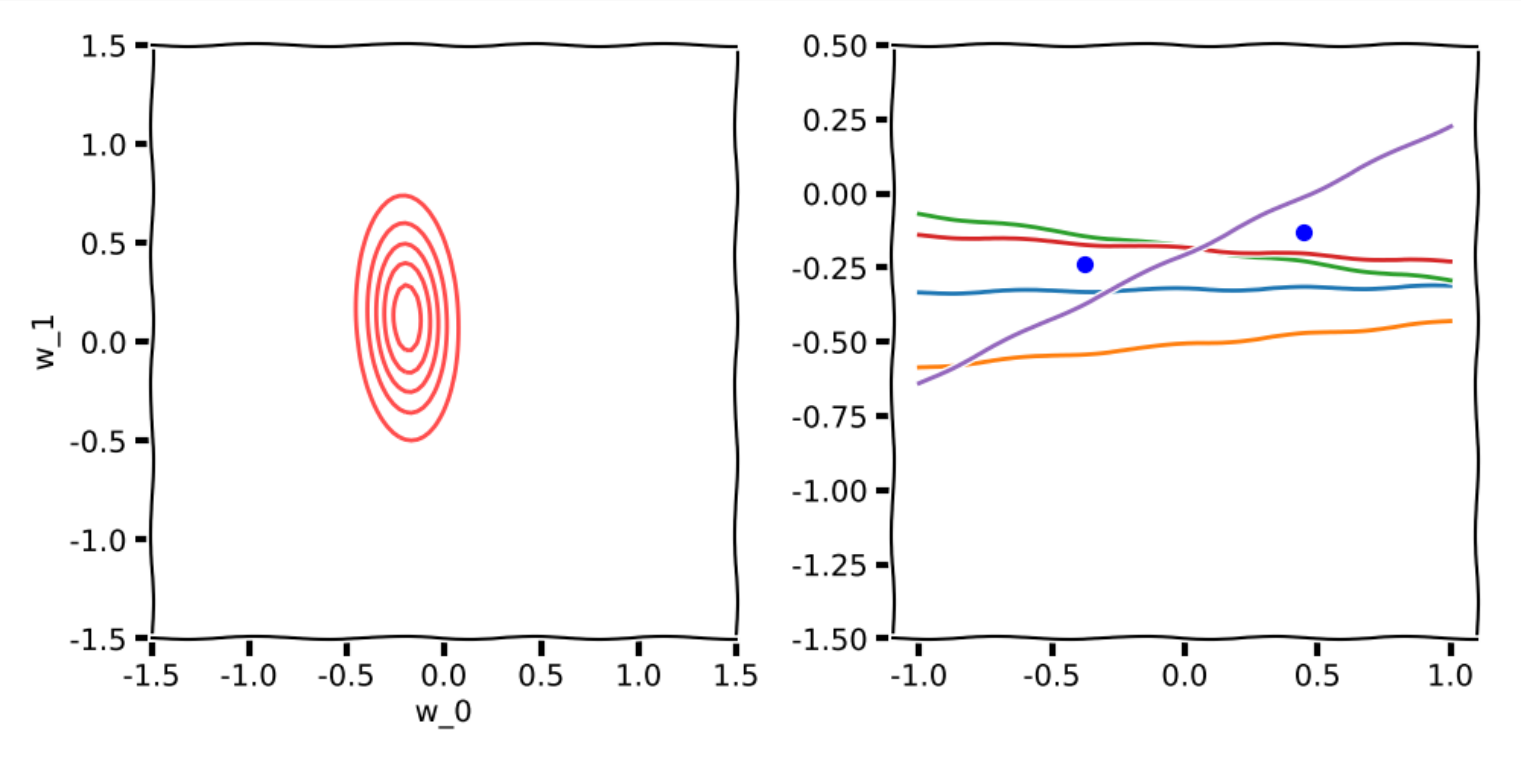

Linear Regression Example

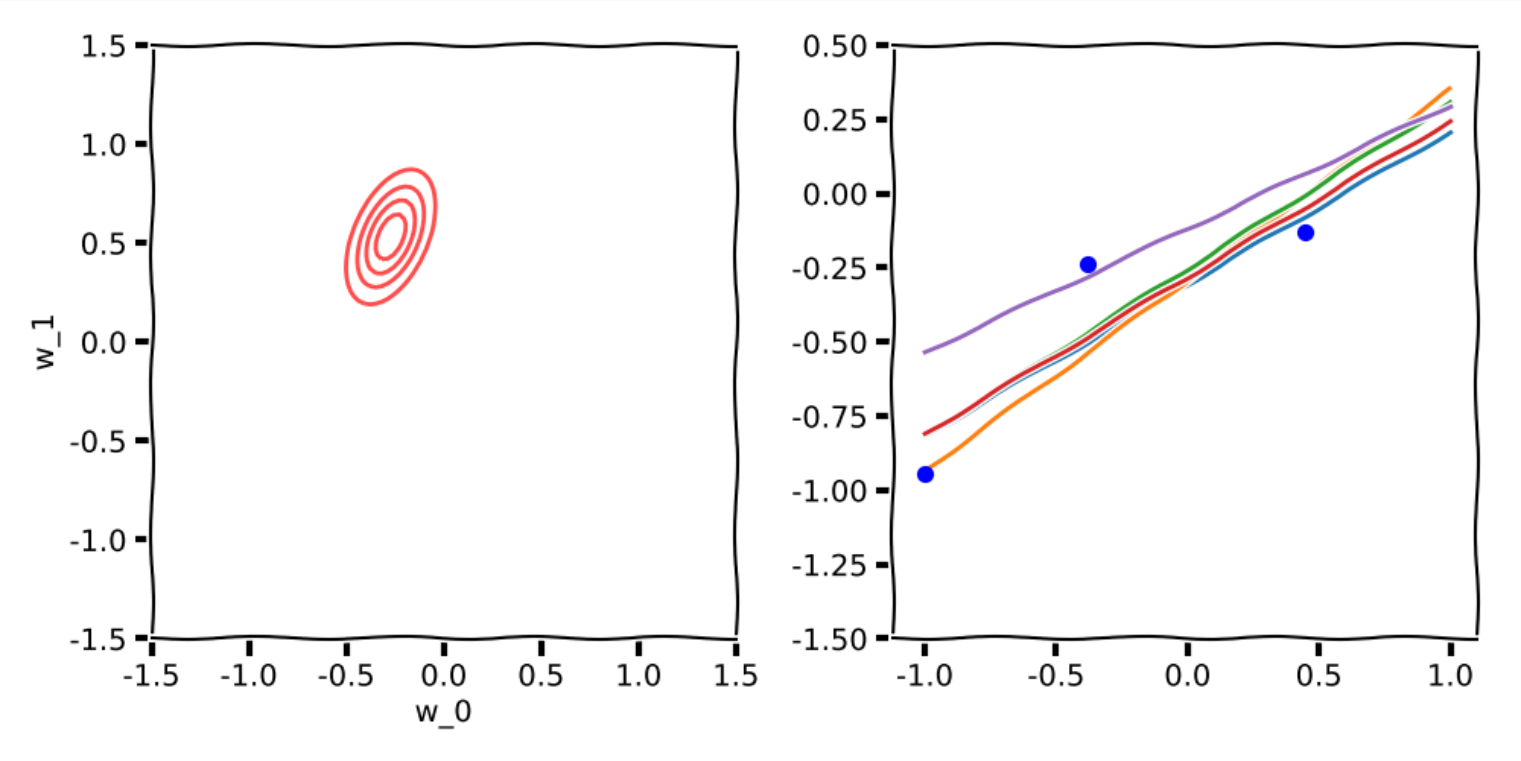

Linear Regression Example

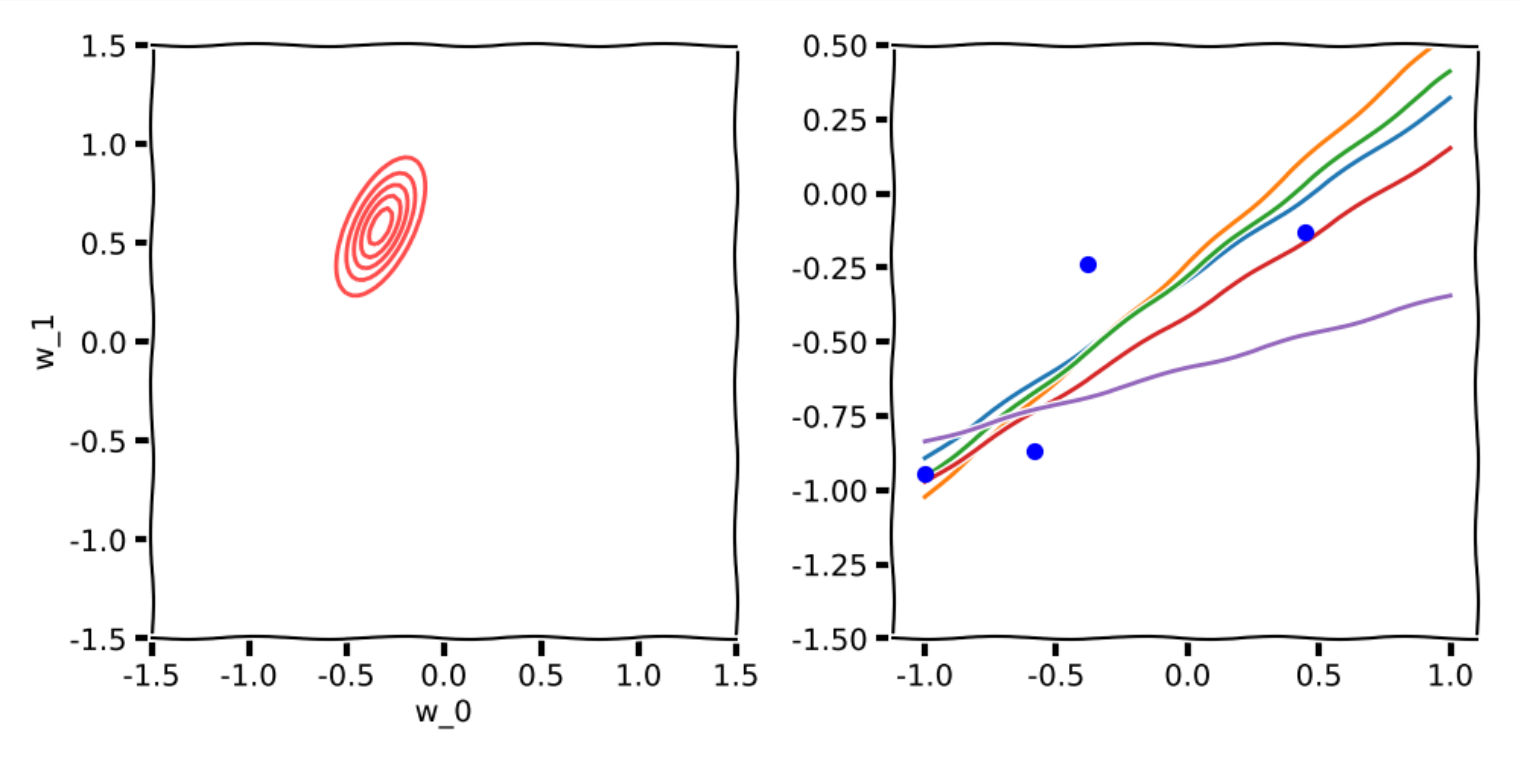

Linear Regression Example

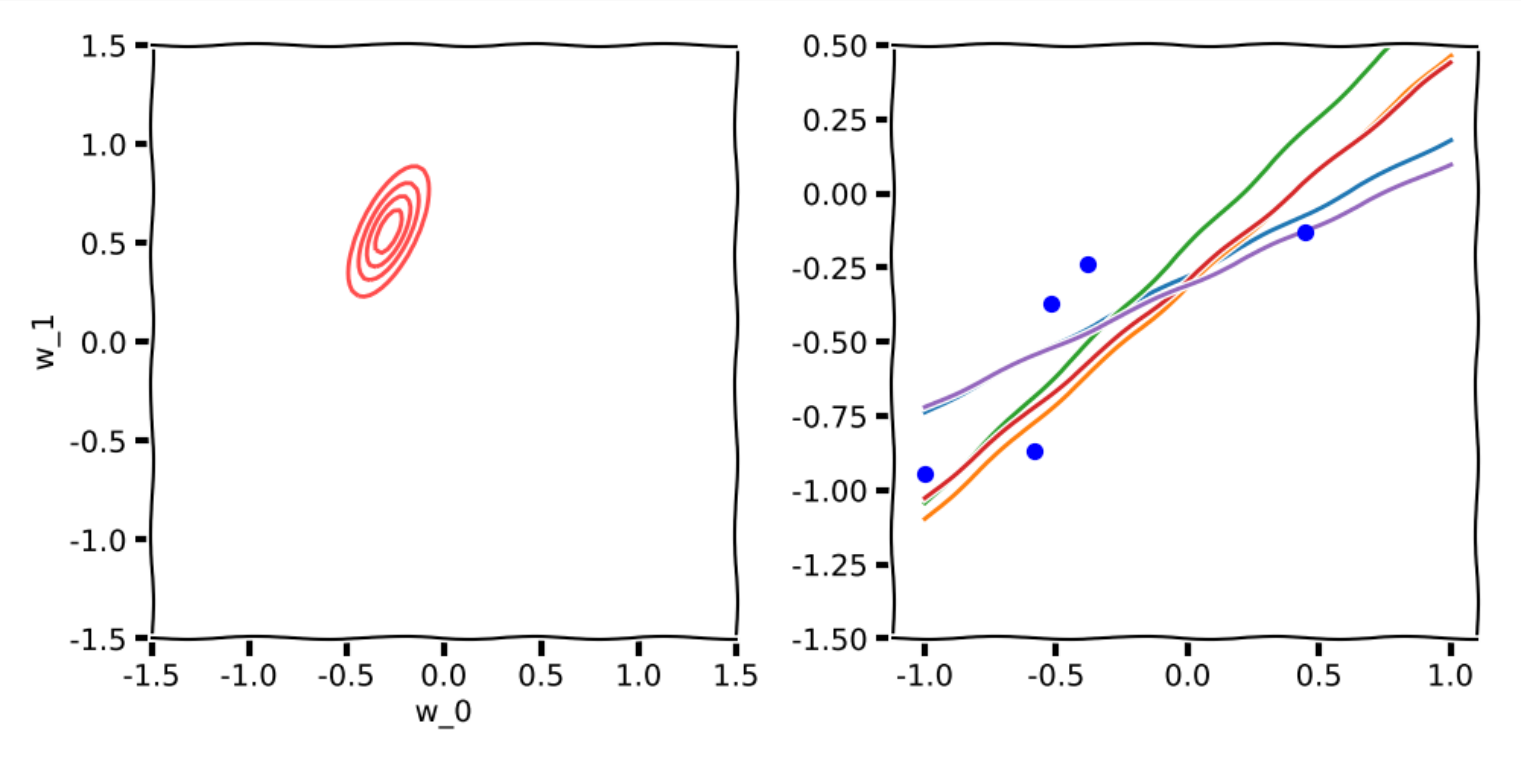

Linear Regression Example

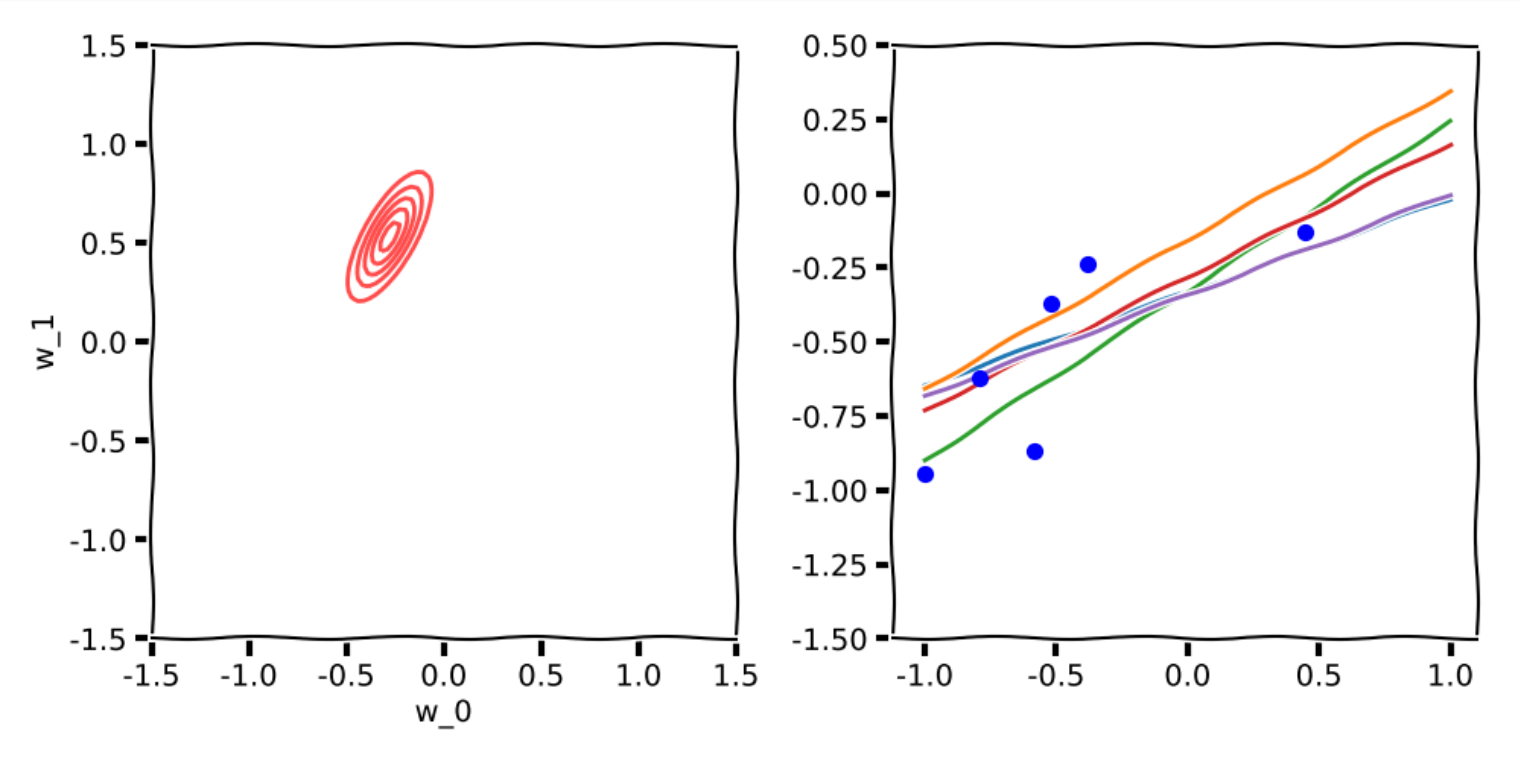

Linear Regression Example

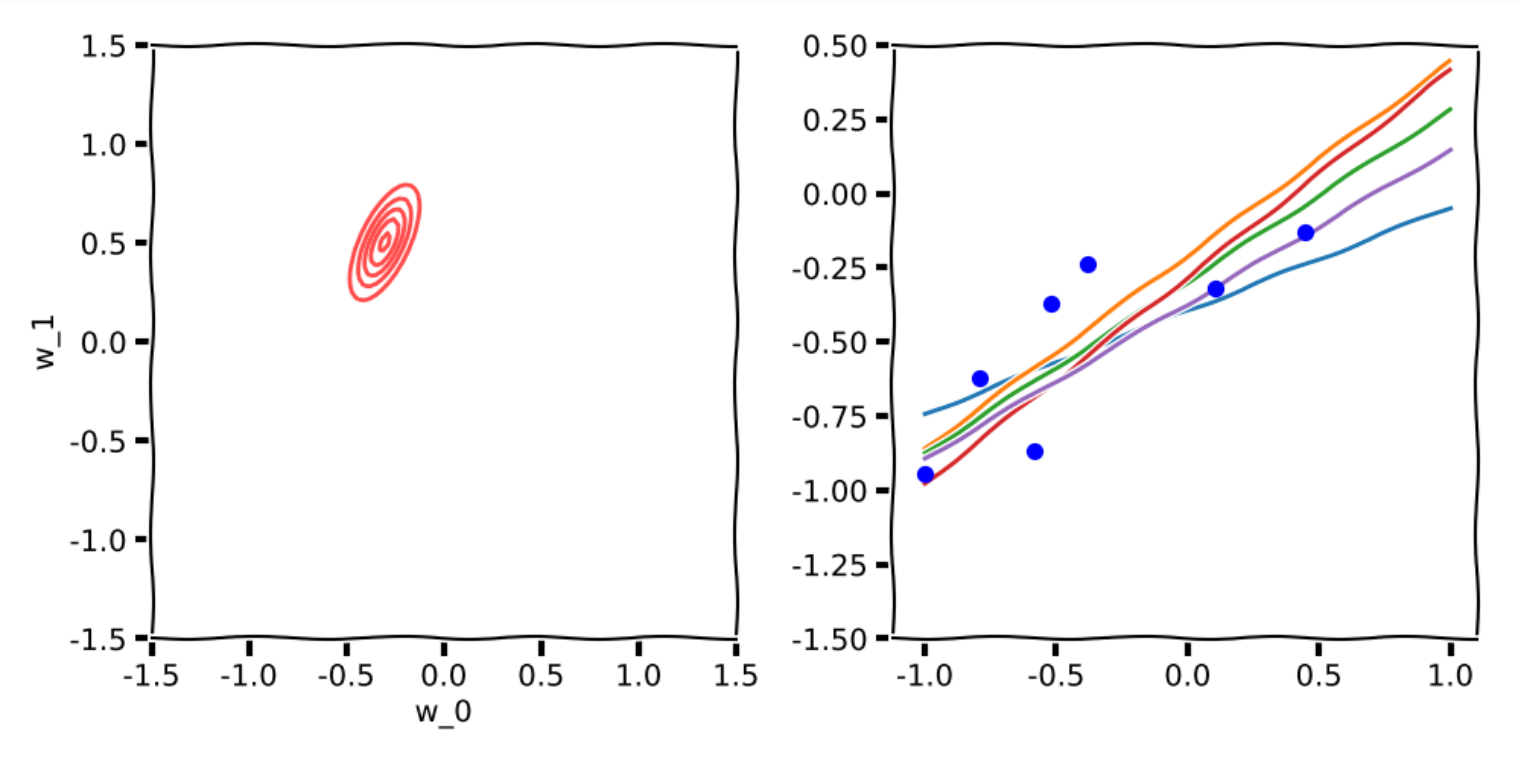

Linear Regression Example

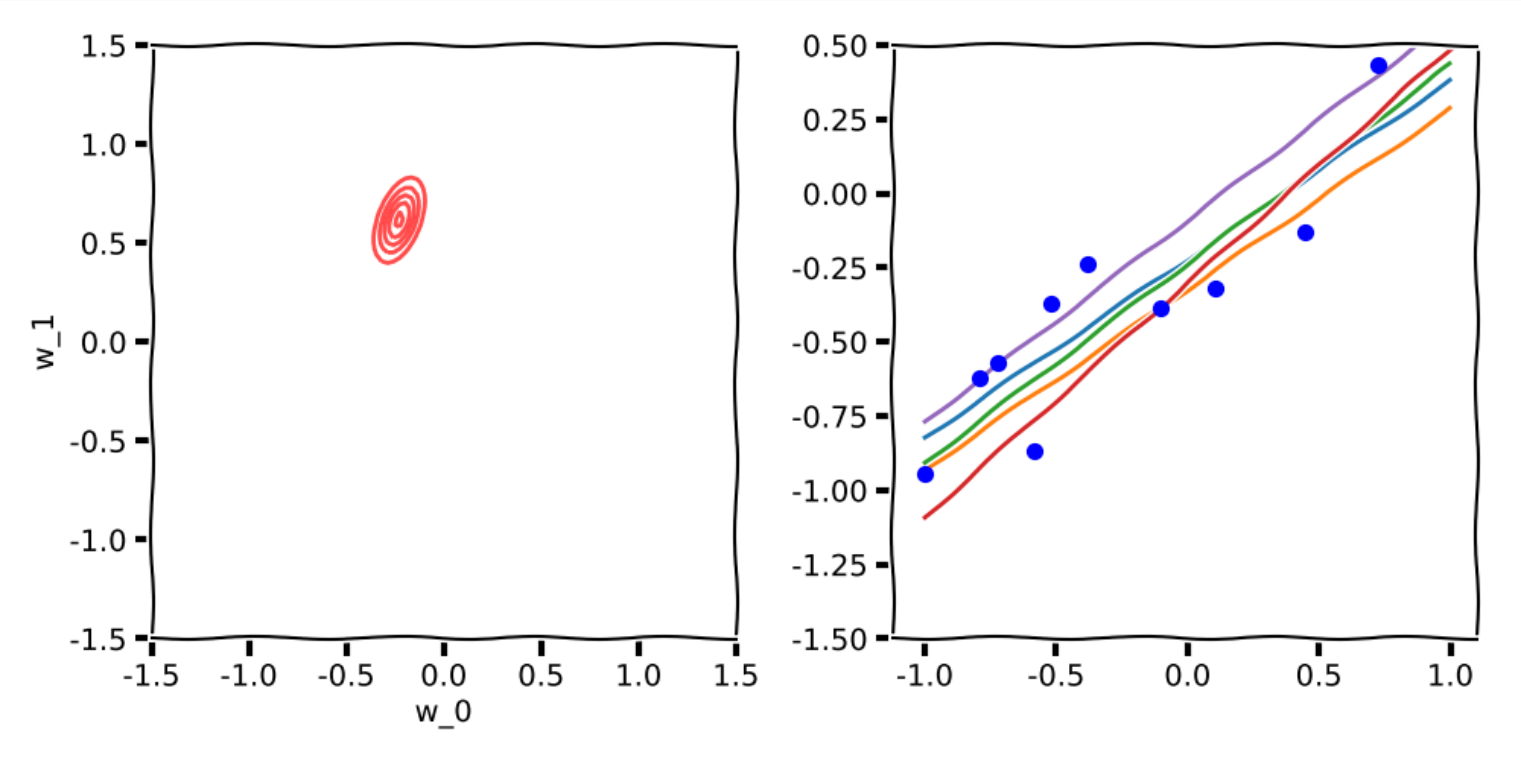

Linear Regression Example

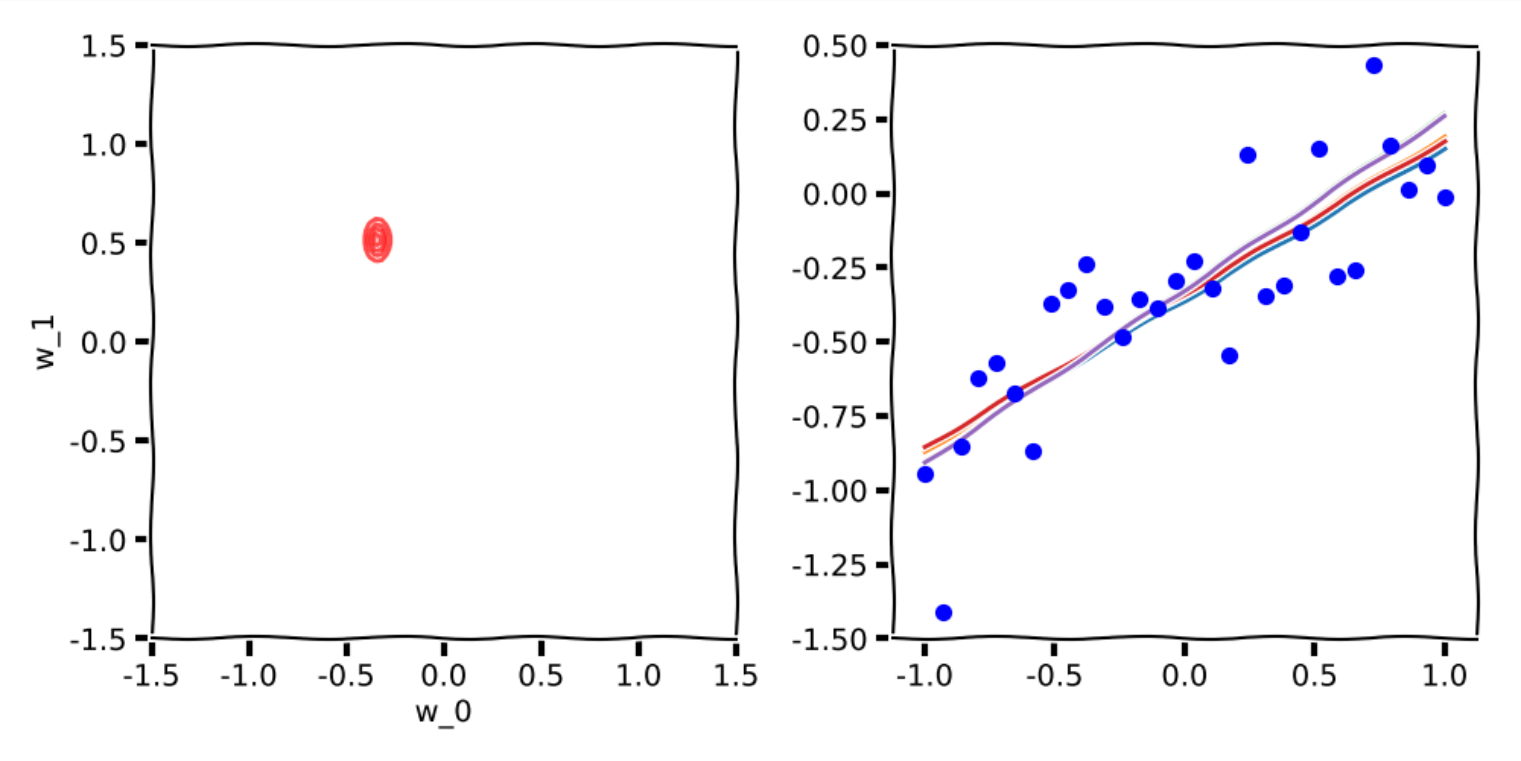

Linear Regression

- Don't underestimate what we just did

Linear Regression

- Don't underestimate what we just did
- We saw data, which we knew where it came from

Linear Regression

- Don't underestimate what we just did
- We saw data, which we knew where it came from
- We made an assumption

Linear Regression

- Don't underestimate what we just did
- We saw data, which we knew where it came from
- We made an assumption
- We recovered the system

Linear Regression

- Don't underestimate what we just did
- We saw data, which we knew where it came from
- We made an assumption
- We recovered the system
- We generated knowledge from data!!!

Linear Regression

- Don't underestimate what we just did
- We saw data, which we knew where it came from
- We made an assumption
- We recovered the system
- We generated knowledge from data!!!
- Understand [1] 3.3

Statistics or Machine Learning

"The difference between statistics and machine learning is that the former cares about parameters while the latter cares about prediction"

– Prof. Neil D. Lawrence

Prediction

$$p(t_*|\mathbf{t}, \mathbf{x}_*, \mathbf{X}, \alpha, \beta) = \int p(t_*|\mathbf{x}_*, \mathbf{w}, \beta) p(\mathbf{w}|\mathbf{t}, \mathbf{X}, \alpha, \beta) \mathrm{d}\mathbf{w}$$

- we do not really care about w we care about new prediction t_* at location $\mathbf{x}_*$
- look at the marginal distribution, i.e. when we average out the weight
- integrate a Gaussian over a Gaussian $\Rightarrow$ Gaussian identities

Prediction

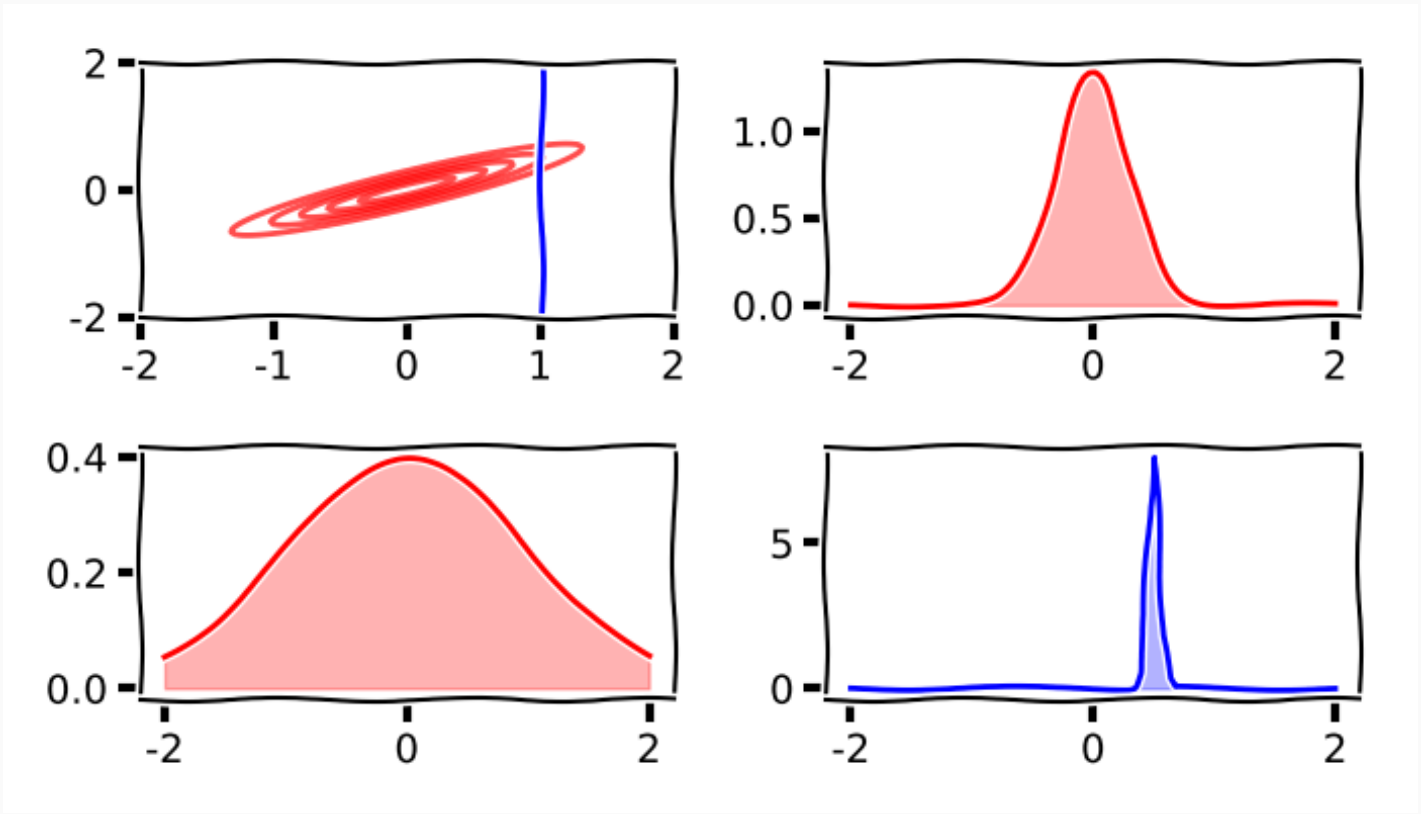

$$p(t_*|\mathbf{t}, \mathbf{x}_*, \mathbf{X}, \alpha, \beta) = \int p(t_*|\mathbf{x}_*, \mathbf{w}, \beta) p(\mathbf{w}|\mathbf{t}, \mathbf{X}, \alpha, \beta) \mathrm{d}\mathbf{w}$$

$$\mathcal{N}(t_*|\mathbf{m}_N^{\mathrm{T}} \phi(\mathbf{x}_*), \frac{1}{\beta} + \phi(\mathbf{x}_*)^{\mathrm{T}} \mathbf{S}_N \phi(\mathbf{x}_*))$$

Predictive Posterior

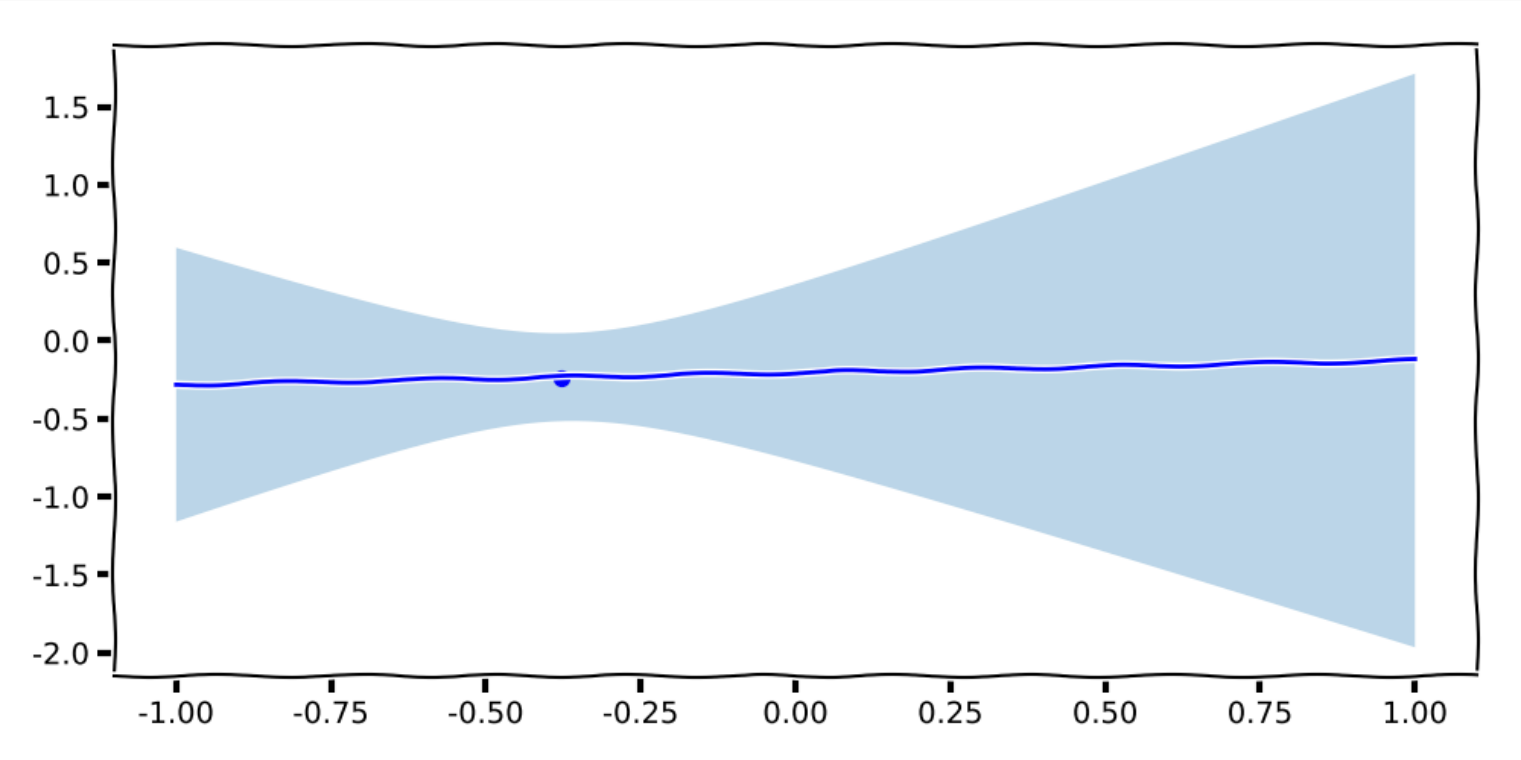

Predictive Posterior

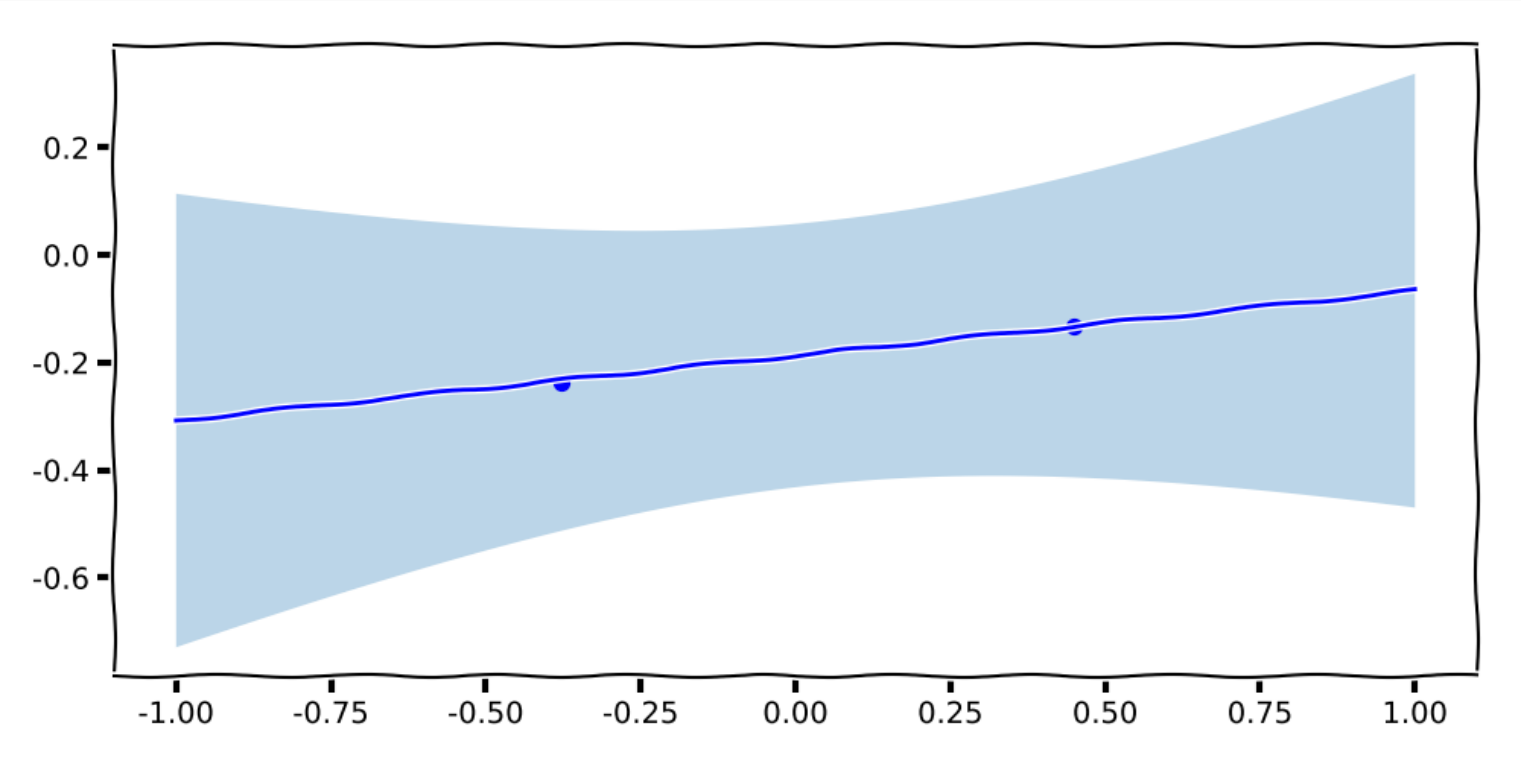

Predictive Posterior

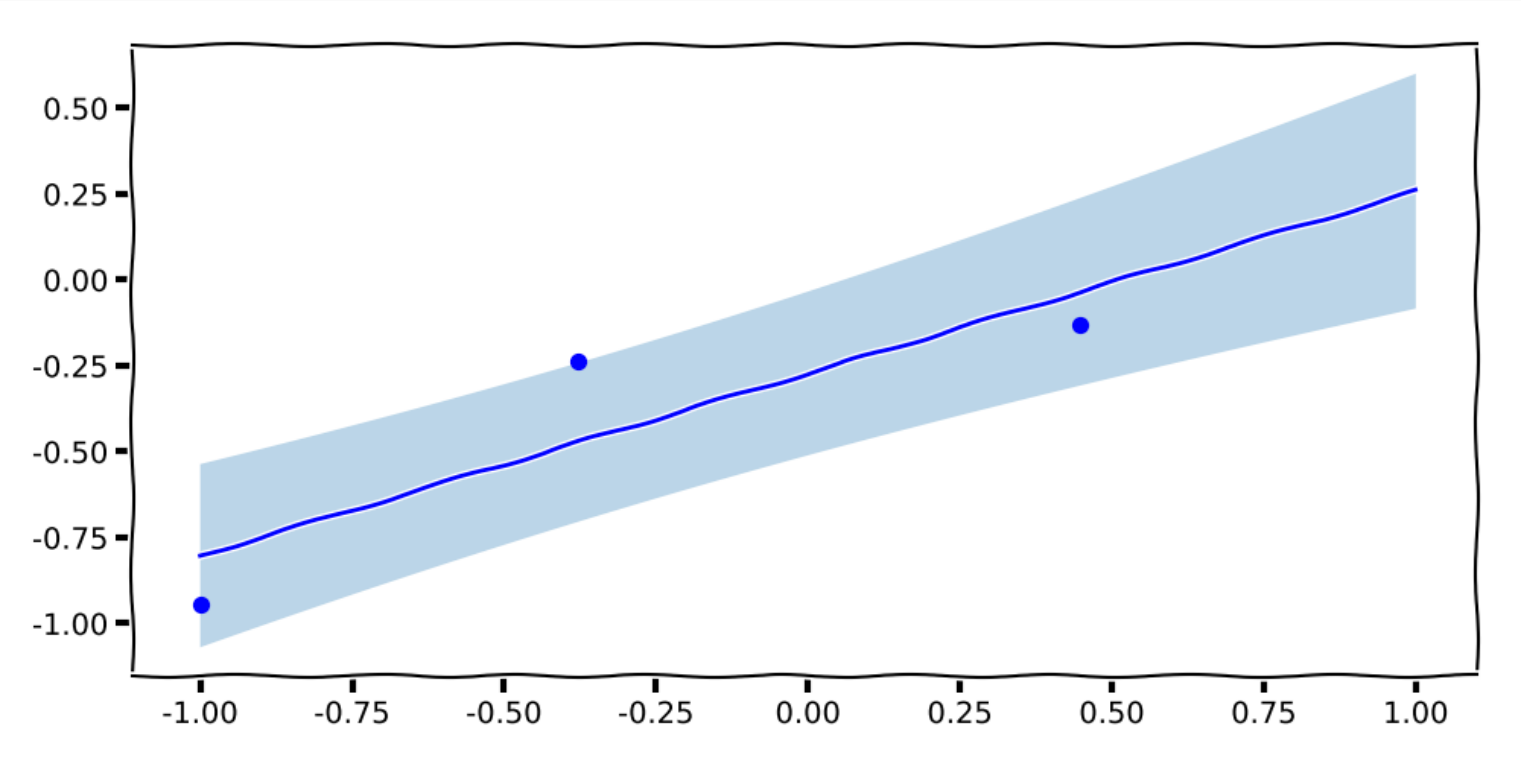

Predictive Posterior

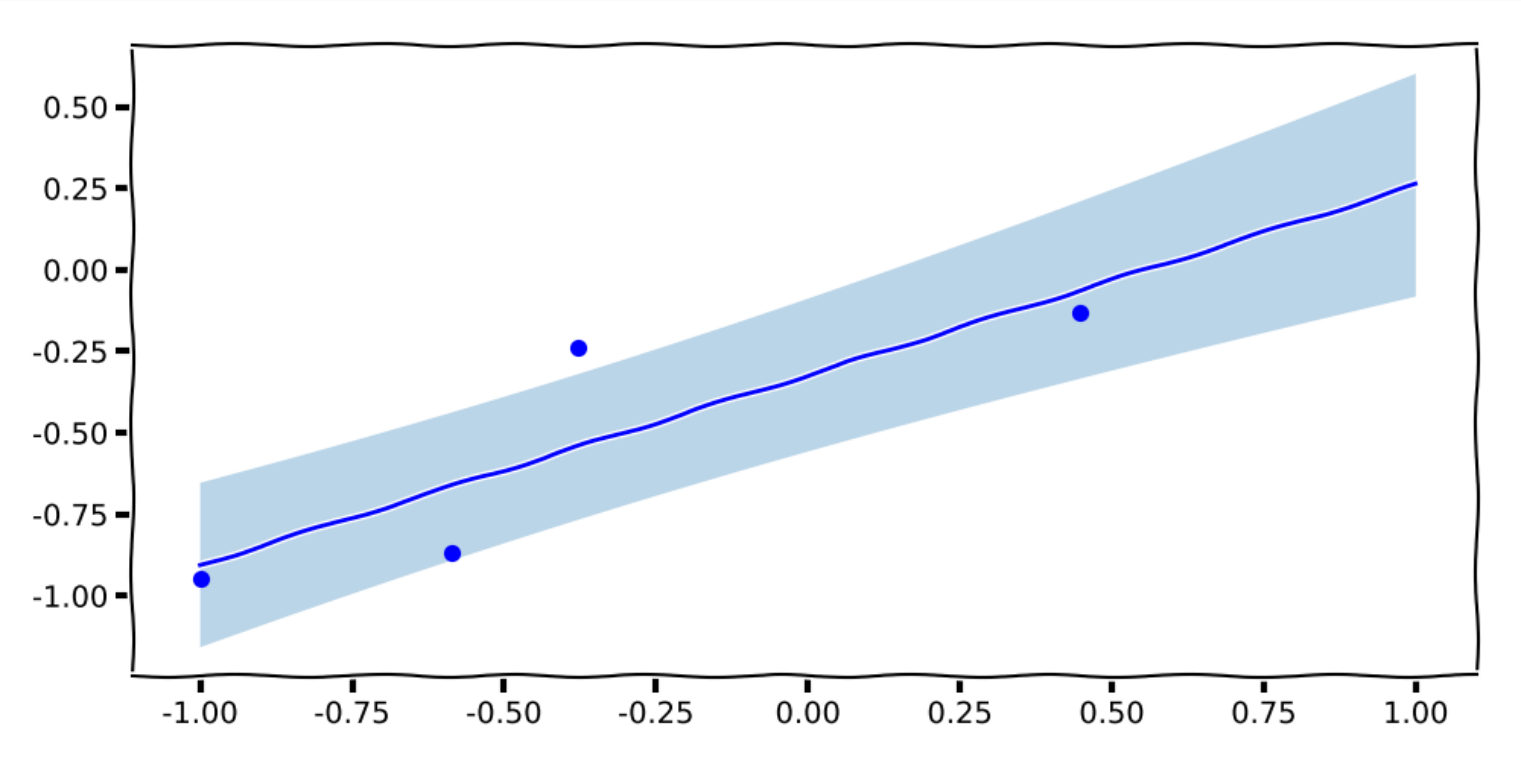

Predictive Posterior

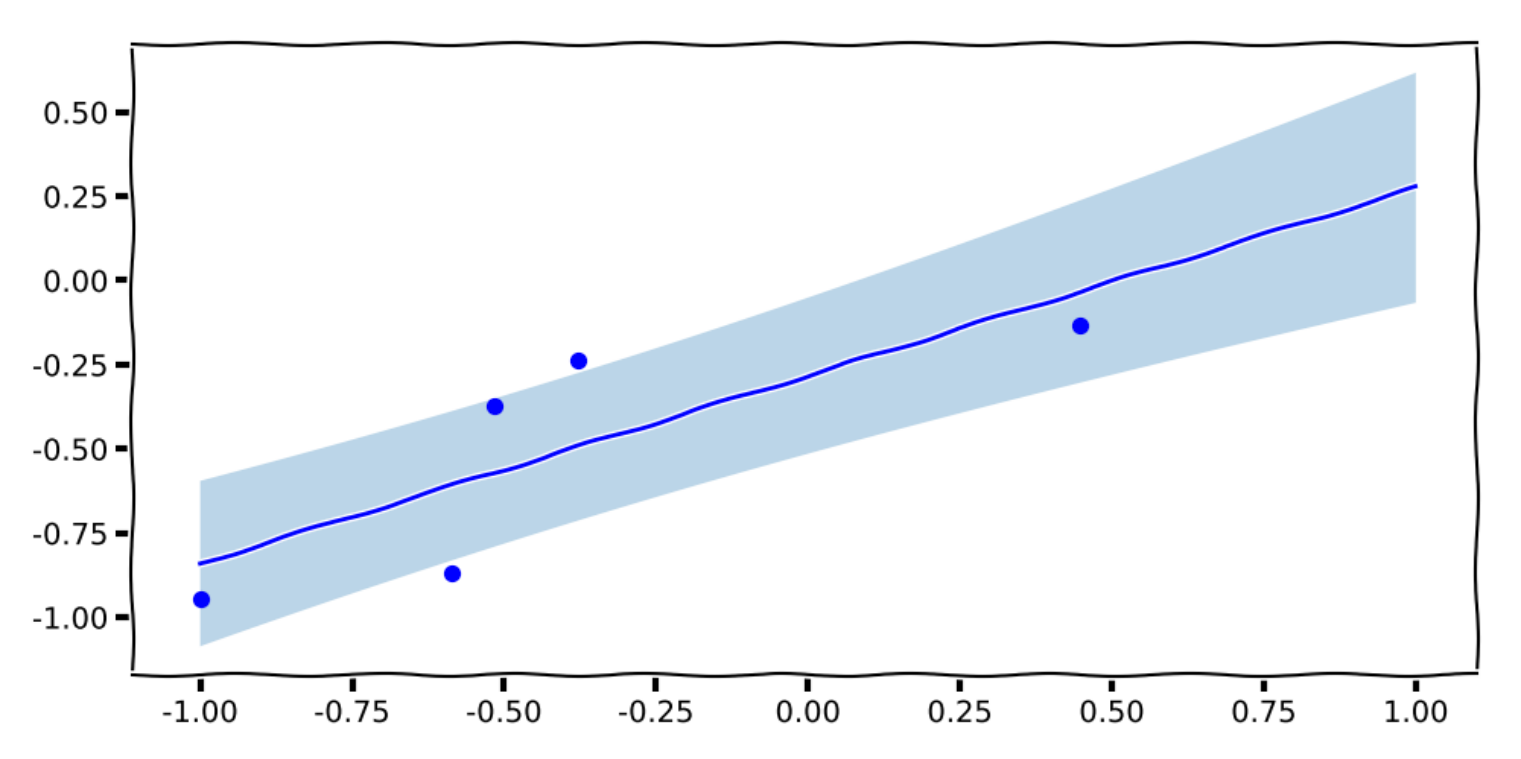

Predictive Posterior

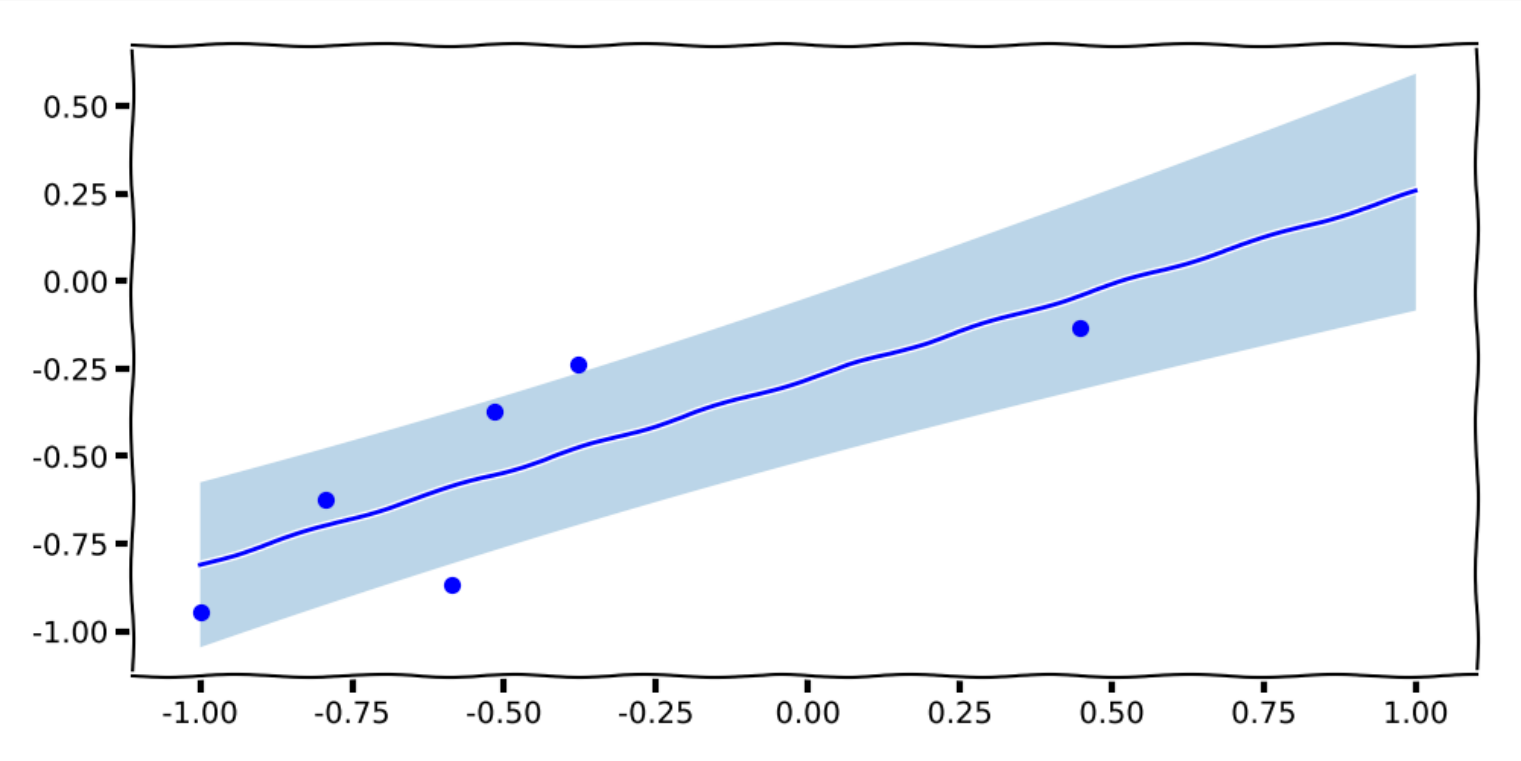

Predictive Posterior

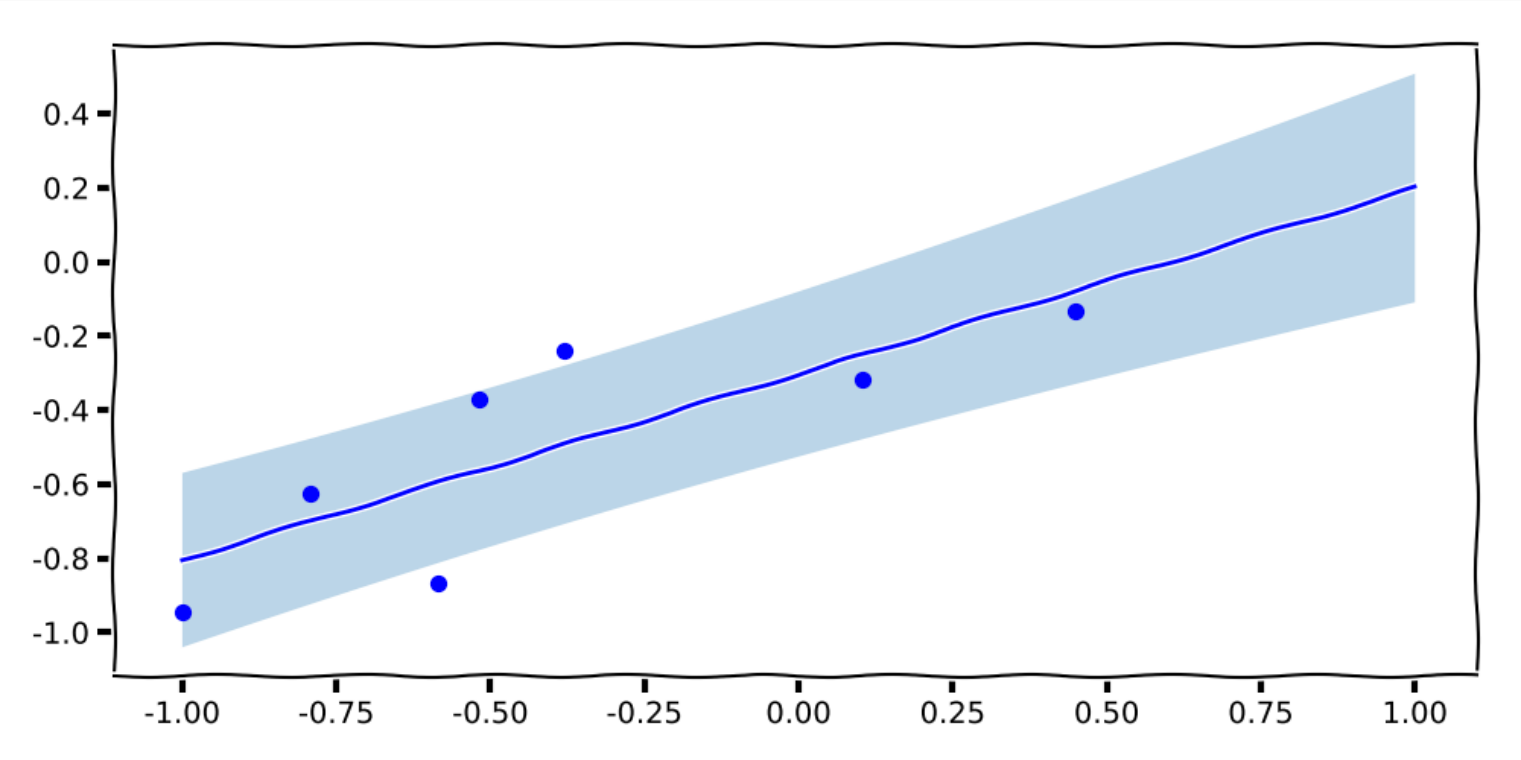

Predictive Posterior

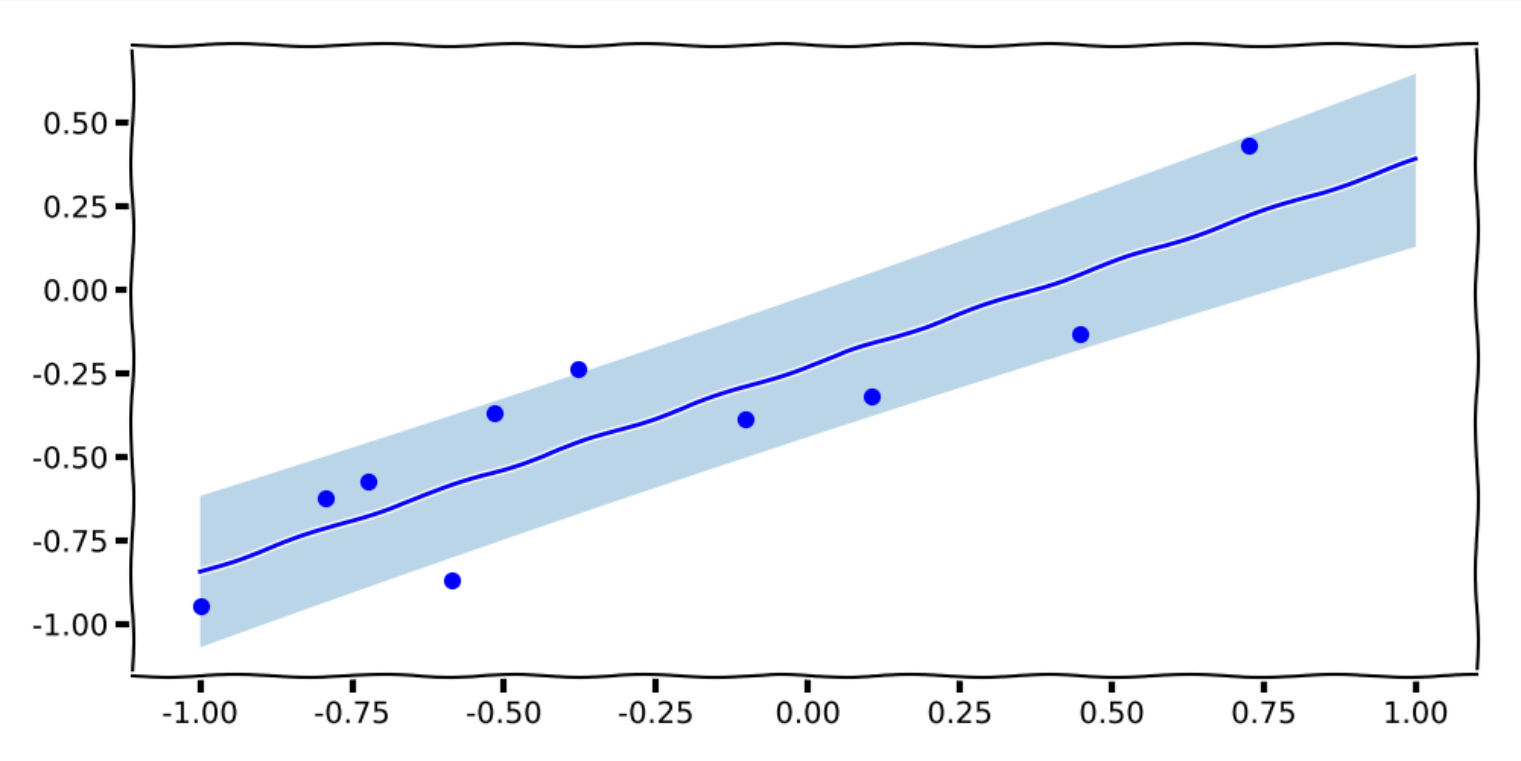

Predictive Posterior

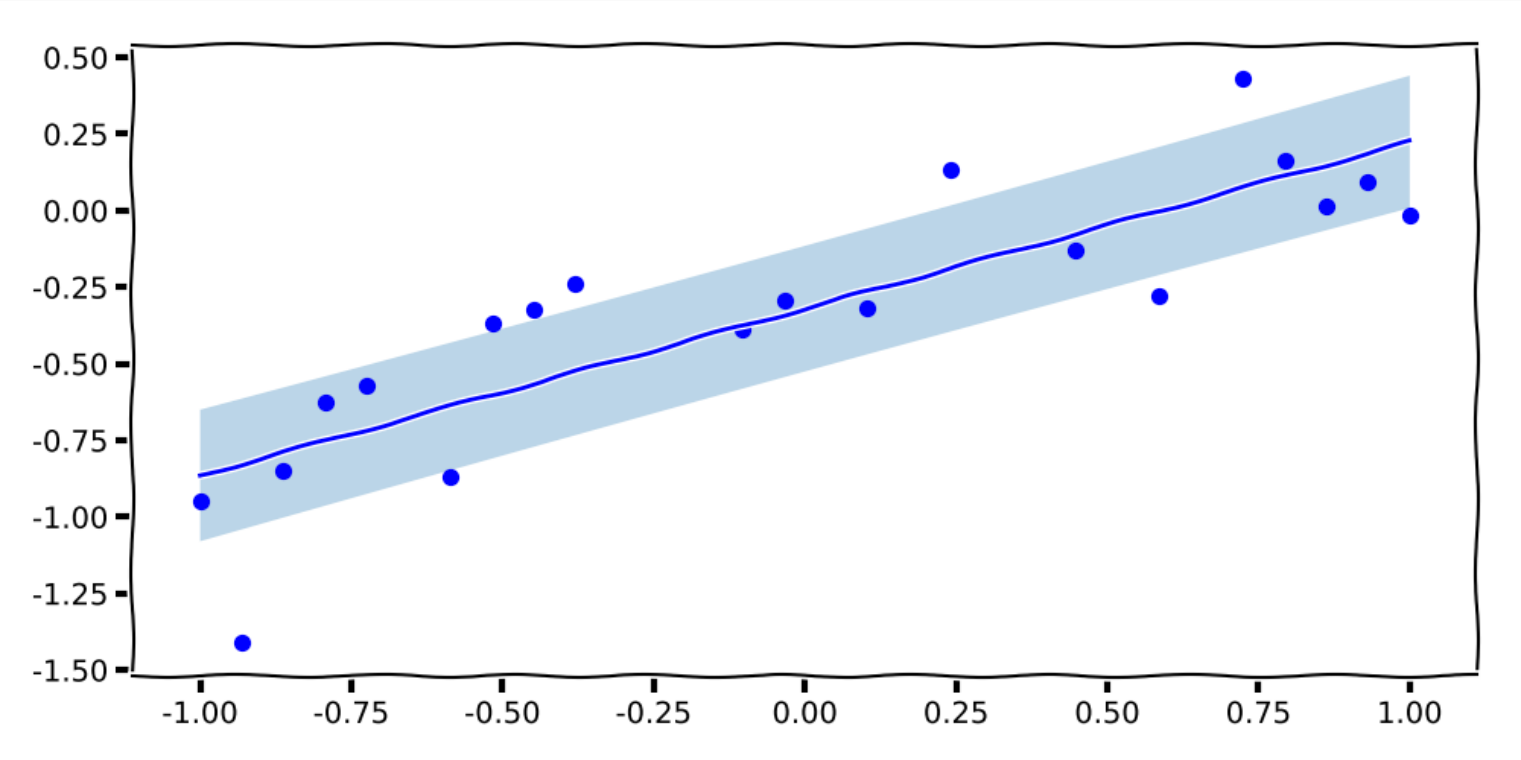

Predictive Posterior

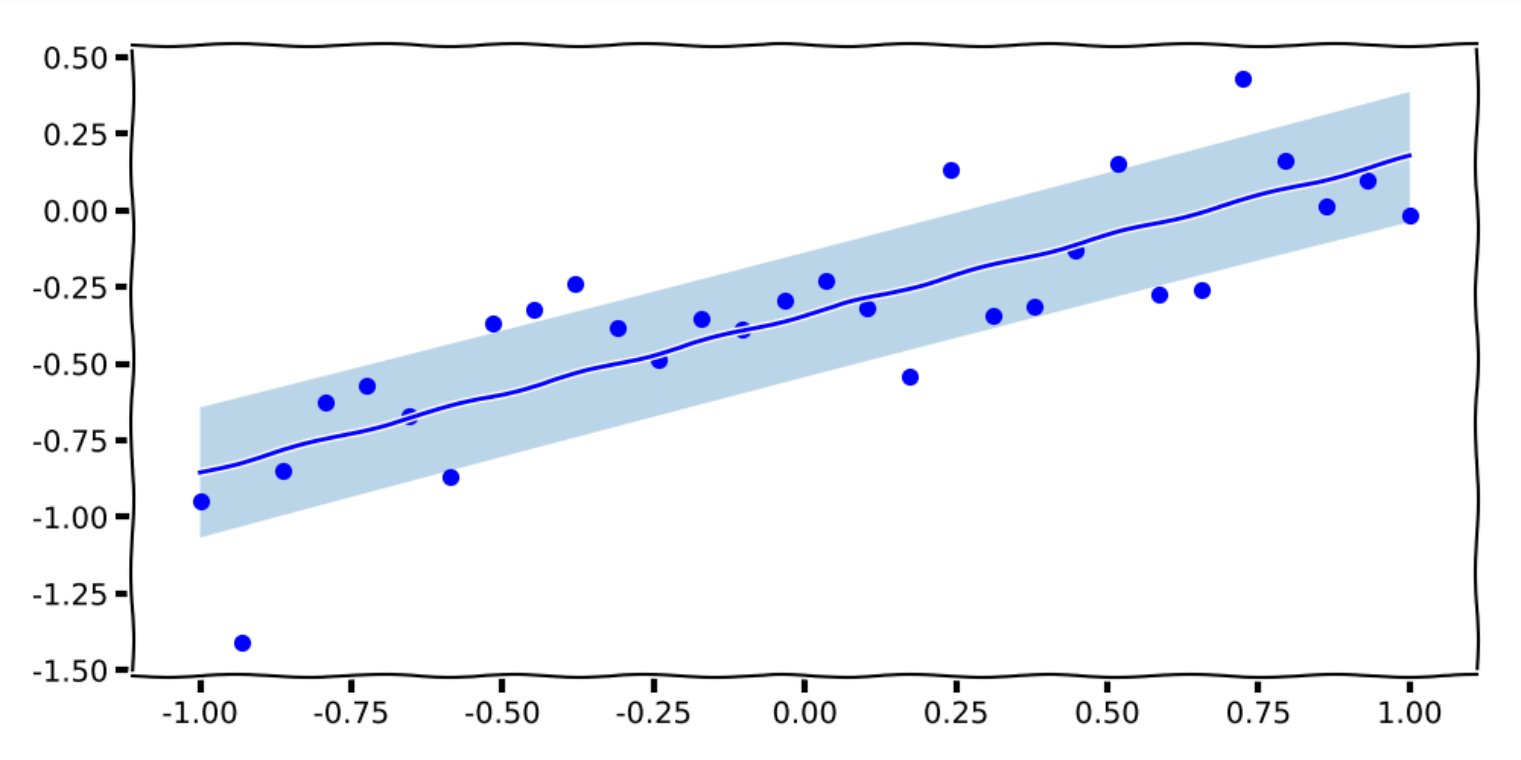

Signal and Noise

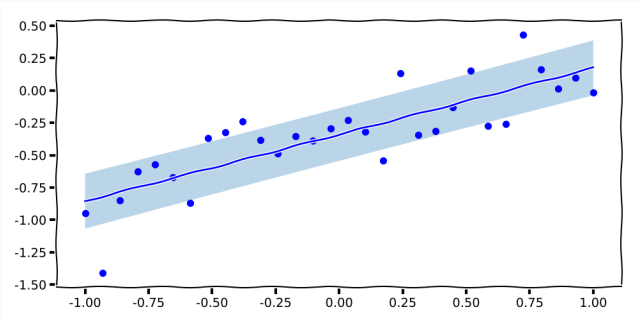

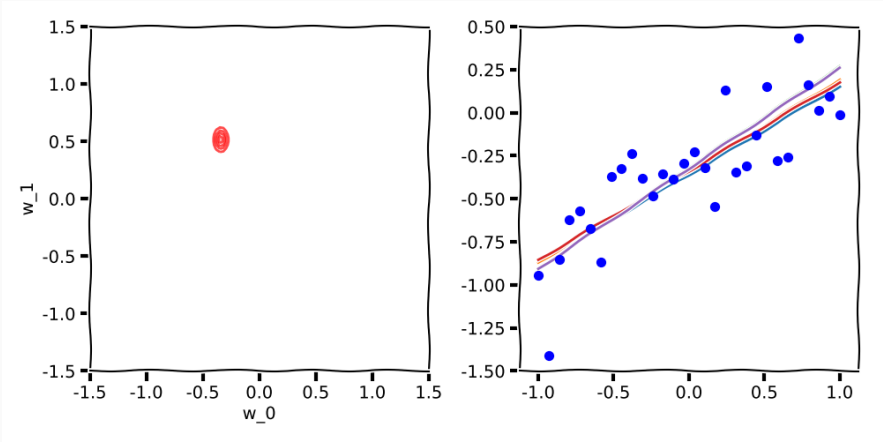

Predictive Posterior [1] Figure 3.8

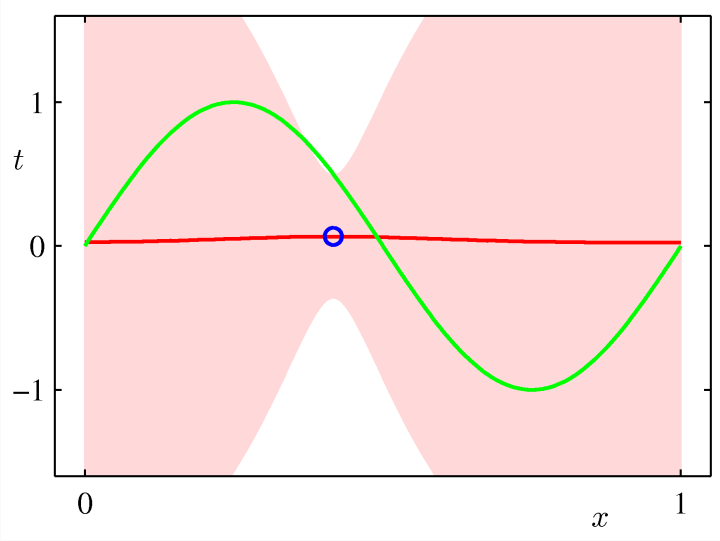

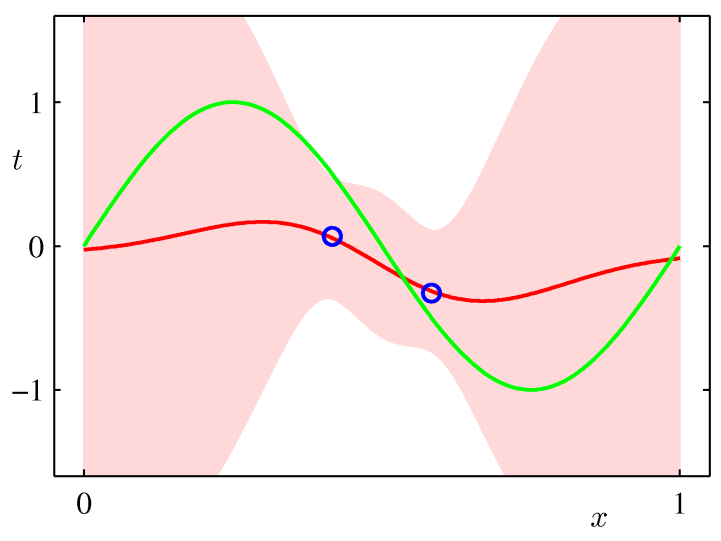

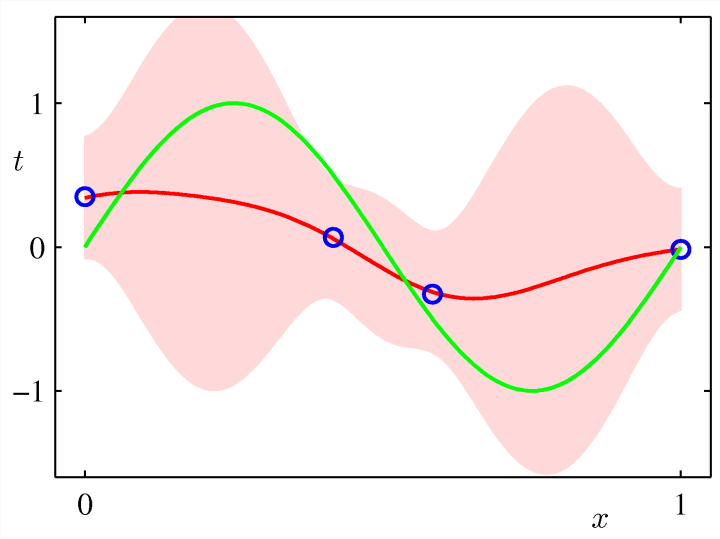

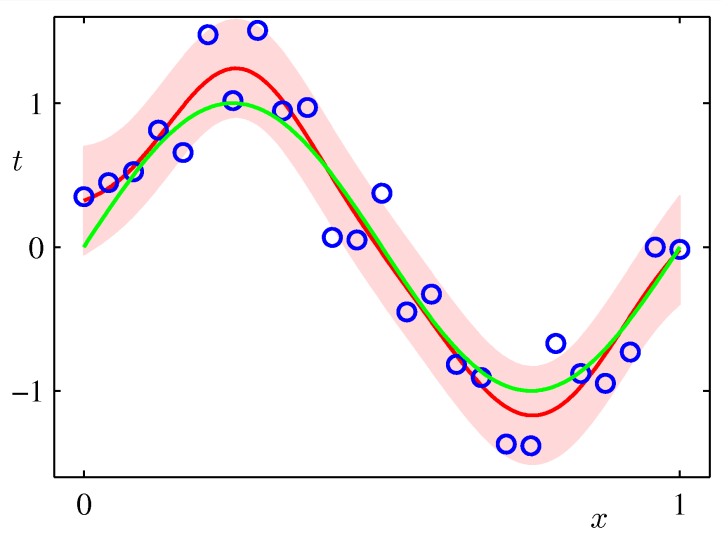

Which Parametrisation

- Should I use a line, polynomial, quadratic basis function?
- Likelihood won't help me
- How do we proceed?

Marginal Distribution

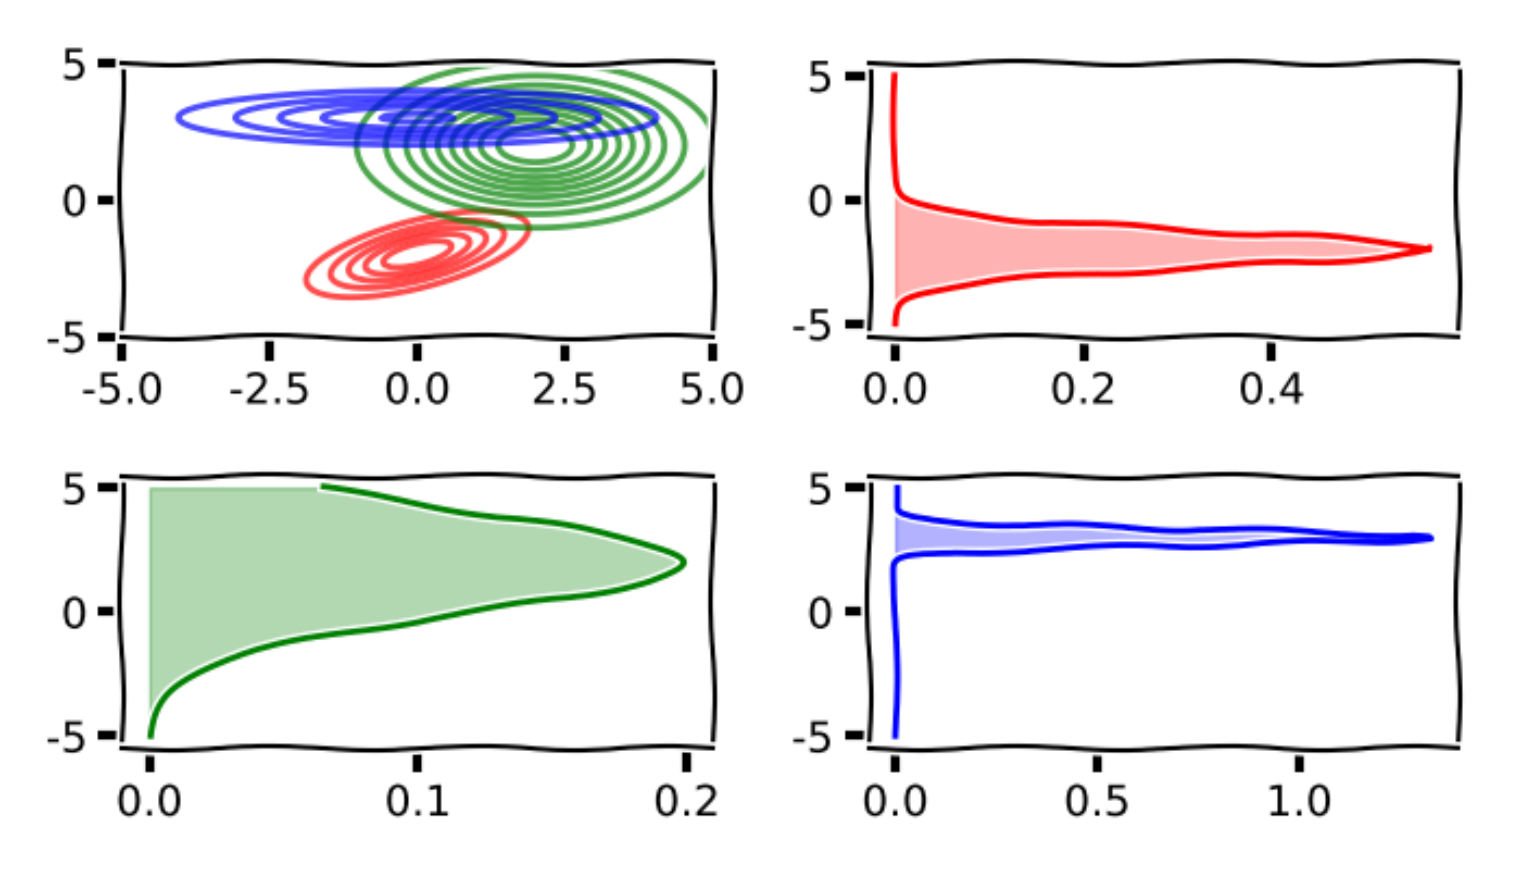

Marginal Likelihood [1] Figure 3.13

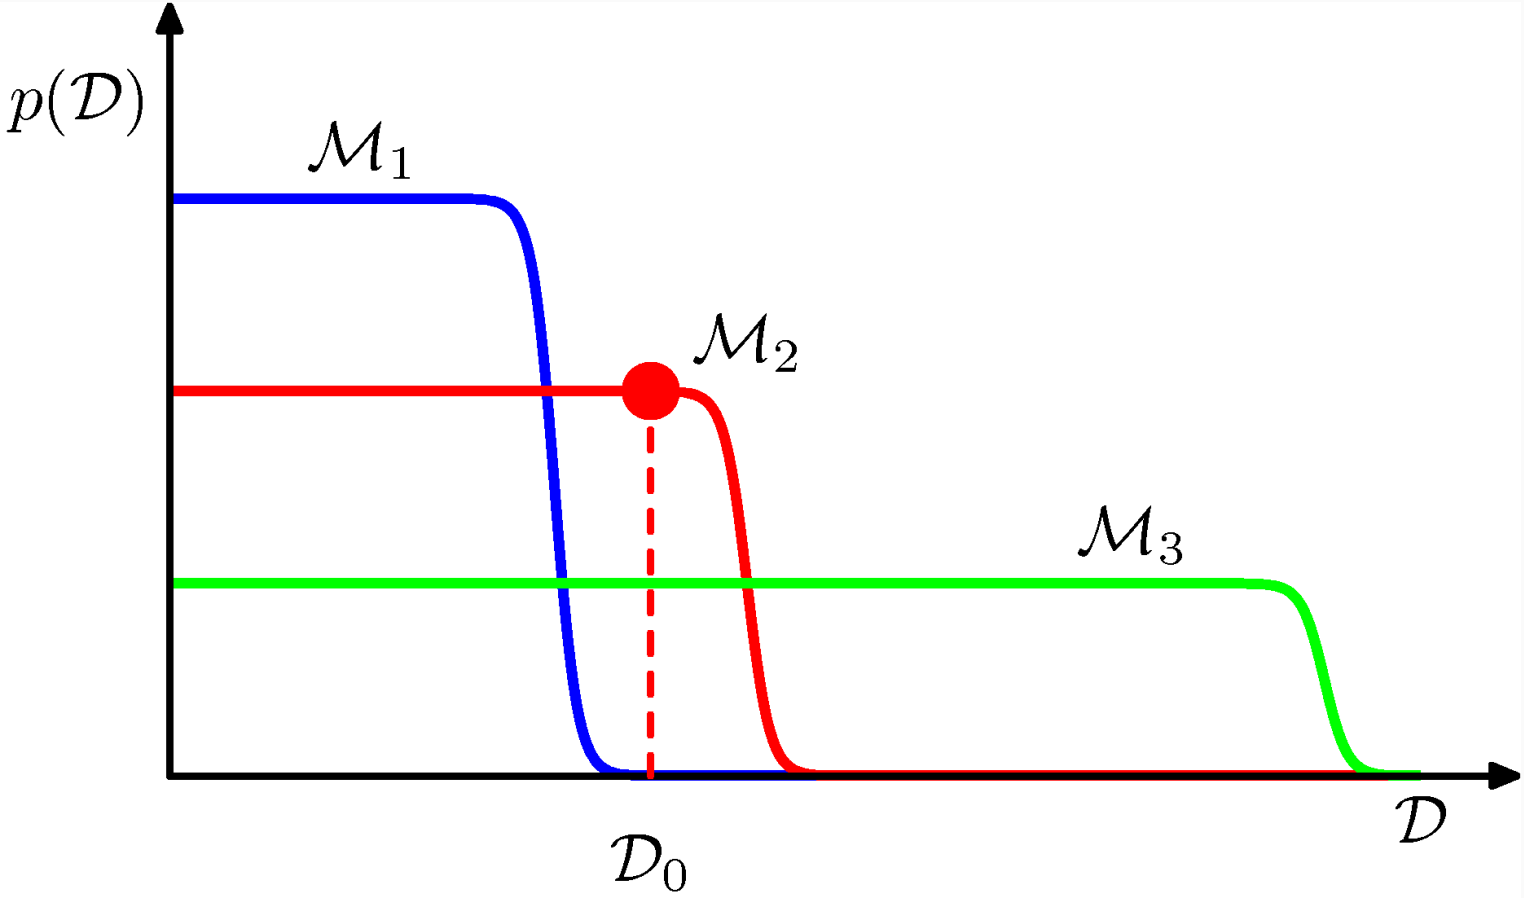

Summary

So Far

Lecture 1 What is machine Learning

- assumptions are the fundation of learning
- probabilities are the language of assumptions

Lecture 2 Probabilities

- what are the rules of probability
- distributions are the parametrised form of a probability

Lecture 3 Distributions

- discrete and continous distributions
- conjugate distributions

Today Models

- how to apply our assumptions to data
- how to learn for real

Next week

- Linear models can only take us that far
    - Monday - Non-linear models
- Fixed model complexity
    - Tuesday - Non-parametric models

Question 1-6 12

References


Christopher M. Bishop.
Pattern Recognition and Machine Learning (Information Science and Statistics).
Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2006.